

# **11-755— Spring 2021**

## **Large Scale Multimedia Processing**



# Lecture 2/6

## Multimedia processing


Rita Singh

Carnegie Mellon University

# In this lecture


- Digital multimedia: Recording and devices
  - Audio
  - Images
  - Video
  - Text
- Digital multimedia: Processing
  - Audio processing
  - Two generic processing techniques

# Mathematical abstraction of sinusoids



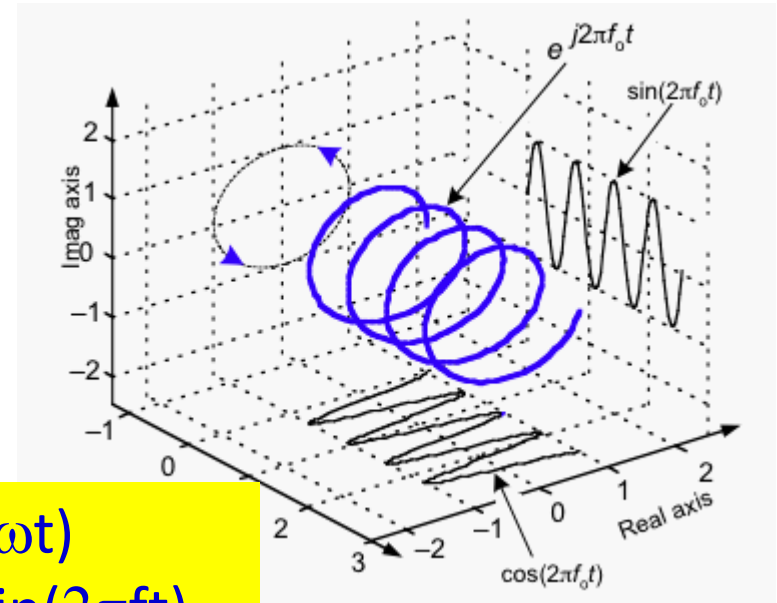
- A phasor abstraction for sinusoids

# Mathematical abstraction of sinusoids




$$e^{j\omega t} = \cos(\omega t) + j \sin(\omega t)$$

$$e^{j2\pi f t} = \cos(2\pi f t) + j \sin(2\pi f t)$$




- $f$ : cycles/sec
- Per cycle:  $2\pi$  radians
- phase = (cycles/sec)  $\times$  (angles/cycle)  $\times$  time =  $f 2\pi t$  (cycles per second as unit)
  - Think of it as total angle traversed over time
- Angle/sec is denoted as  $\omega = f 2\pi$  (also called angular frequency)
- Phase -  $\omega t$  (angles per second as unit)

# Signal decomposition



- Any periodic signal can be represented as a sum of sinusoids

# The complex exponential



- The complex exponential is a complex sum of two sinusoids
  - $e^{j\theta} = \cos\theta + j \sin\theta$
- The real part is a cosine function
- The imaginary part is a sine function
- A complex exponential time series is a complex sum of two time series
  - $e^{j\omega t} = \cos(\omega t) + j \sin(\omega t)$
- Two complex exponentials of different frequencies are “orthogonal” to each other. i.e.


$$\int_{-\infty}^{\infty} e^{j\alpha t} e^{j\beta t} dt = 0 \quad \text{if } \alpha \neq \beta$$

# The discrete Fourier transform


- The discrete Fourier transform decomposes the signal into the sum of a finite number of complex exponentials
  - As many exponentials as there are samples in the signal being analyzed
- An aperiodic signal *cannot* be decomposed into a sum of a finite number of complex exponentials
  - Or into a sum of any countable set of periodic signals
- The discrete Fourier transform actually assumes that the signal being analyzed is exactly one period of an infinitely long signal
  - In reality, it computes the Fourier spectrum of the infinitely long periodic signal, of which the analyzed data are one period



# The discrete Fourier transform



- The discrete Fourier transform of the above signal **actually computes the Fourier spectrum of the periodic signal** shown below
  - Which extends from  $-\infty$  to  $+\infty$
  - The period of this signal is 31 samples in this example



# The discrete Fourier transform

- The  $k^{\text{th}}$  point of a Fourier transform is computed as:

$$X[k] = \sum_{n=0}^{M-1} x[n] e^{-\frac{j2\pi kn}{M}}$$

- $x[n]$  is the  $n^{\text{th}}$  point in the analyzed data sequence
  - $X[k]$  is the value of the  $k^{\text{th}}$  point in its Fourier spectrum
  - $M$  is the total number of points in the sequence
- Note that the  $(M+k)^{\text{th}}$  Fourier coefficient is identical to the  $k^{\text{th}}$  Fourier coefficient

$$\begin{aligned} X[M+k] &= \sum_{n=0}^{M-1} x[n] e^{-\frac{j2\pi(M+k)n}{M}} = \sum_{n=0}^{M-1} x[n] e^{-\frac{j2\pi Mn}{M}} e^{-\frac{j2\pi kn}{M}} \\ &= \sum_{n=0}^{M-1} x[n] e^{-j2\pi n} e^{-\frac{j2\pi kn}{M}} = \sum_{n=0}^{M-1} x[n] e^{-\frac{j2\pi kn}{M}} = X[k] \end{aligned}$$

# The discrete Fourier transform

- Discrete Fourier transform coefficients are generally complex
  - $e^{j\theta}$  has a real part  $\cos\theta$  and an imaginary part  $\sin\theta$

$$e^{j\theta} = \cos\theta + j \sin\theta$$

- As a result, every  $X[k]$  has the form

$$X[k] = X_{\text{real}}[k] + jX_{\text{imaginary}}[k]$$

- A **magnitude spectrum** represents only the magnitude of the Fourier coefficients

$$X_{\text{magnitude}}[k] = \text{sqrt}(X_{\text{real}}[k]^2 + X_{\text{imag}}[k]^2)$$


- A **power spectrum** is the square of the magnitude spectrum

$$X_{\text{power}}[k] = X_{\text{real}}[k]^2 + X_{\text{imag}}[k]^2$$

- For speech recognition and other audio analyses, we usually use the magnitude or power spectra

# The phasor, magnitude and power

- Magnitude is the radius
  - always positive
- Power is proportional to energy
  - Squared Fourier coefficient
  - always positive




# The discrete Fourier transform


- A discrete Fourier transform of an  $M$ -point sequence will only compute  $M$  unique frequency components
  - i.e. the DFT of an  $M$  point sequence will have  $M$  points
  - The  $M$ -point DFT represents frequencies in the continuous-time signal that was digitized to obtain the digital signal
- The  $0^{\text{th}}$  point in the DFT represents 0Hz, or the DC component of the signal
- The  $(M-1)^{\text{th}}$  point in the DFT represents  $(M-1)/M$  times the sampling frequency
- All DFT points are uniformly spaced on the frequency axis between 0 and the sampling frequency

# The discrete Fourier transform

- A 50 point segment of a decaying sine wave sampled at 8000 Hz



The corresponding 50 point magnitude DFT. The 51<sup>st</sup> point (shown in red) is identical to the 1<sup>st</sup> point.



Sample 0 = 0 Hz

Sample 50 is the 51<sup>st</sup> point  
It is identical to Sample 0

Sample 50 = 8000Hz



# The discrete Fourier transform

- The *Fast Fourier Transform* (FFT) is simply a fast algorithm to compute the DFT
  - It utilizes symmetry in the DFT computation to reduce the total number of arithmetic operations greatly
- The time domain signal can be recovered from its DFT as:

$$x[n] = \frac{1}{M} \sum_{k=0}^{M-1} X[k] e^{j \frac{2\pi kn}{M}}$$

# Example: An audio signal

Tom Sullivan spells out his last name: S U L L I V A N






# Example: A speech signal

Sound is produced as an analog signal.

It is converted to digital format by analog-to-digital (A/D) converters.



This signal has been digitized at the rate of 16000 numbers per second. ie,

Signal sampling frequency=16 kiloHertz

or


Sampling rate = 16 kiloHertz

Samples are 16bit integers. Sample No. 13172:13200 (read left-to-right)

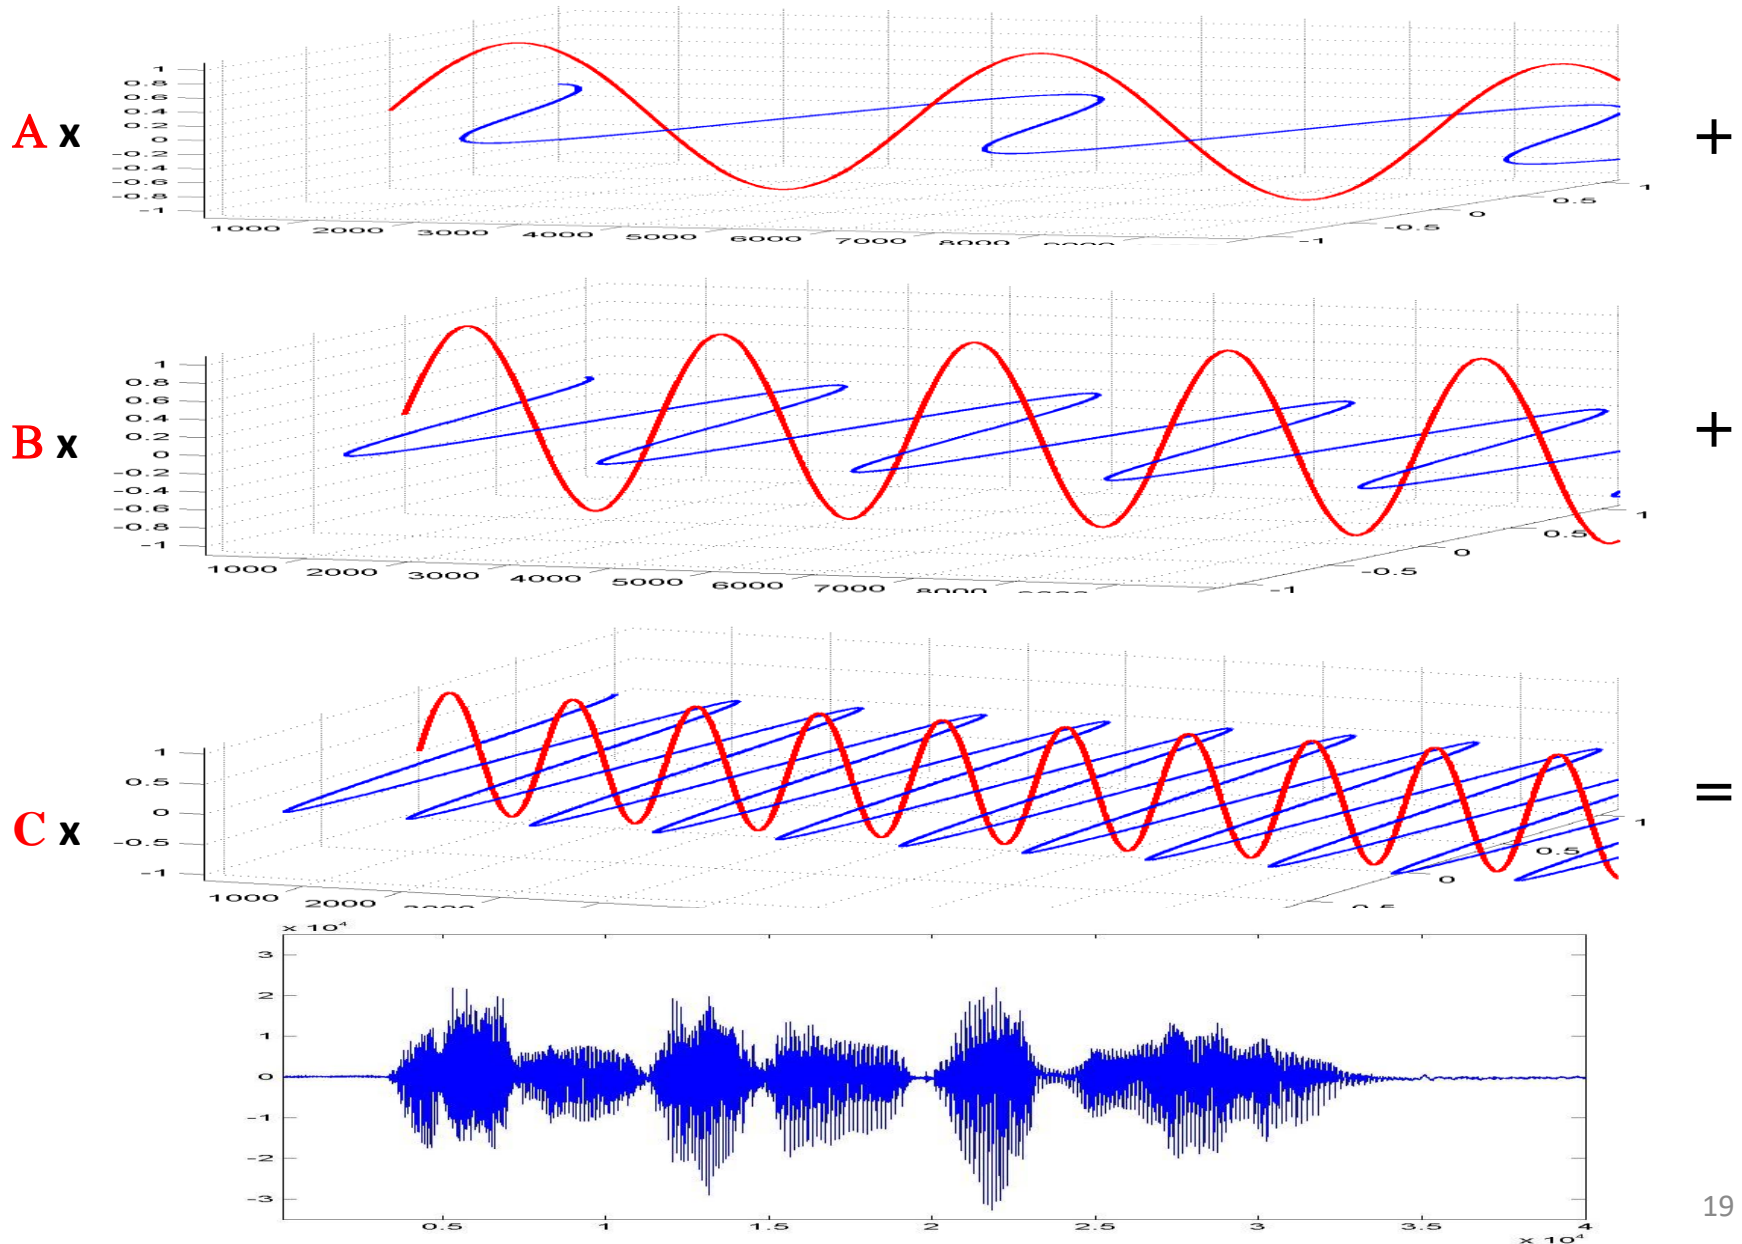
-24780	-29075	-26479	-19223	-18953	-13313	-2011	3440	7412	11209
15823	19792	17963	15337	12193	8742	7267	3721	1150	-443
-2989	-2499	-2761	-4020	-2898	-4016	-6988	-9668	-10614	

# Capturing the Spectrum: The discrete Fourier transform

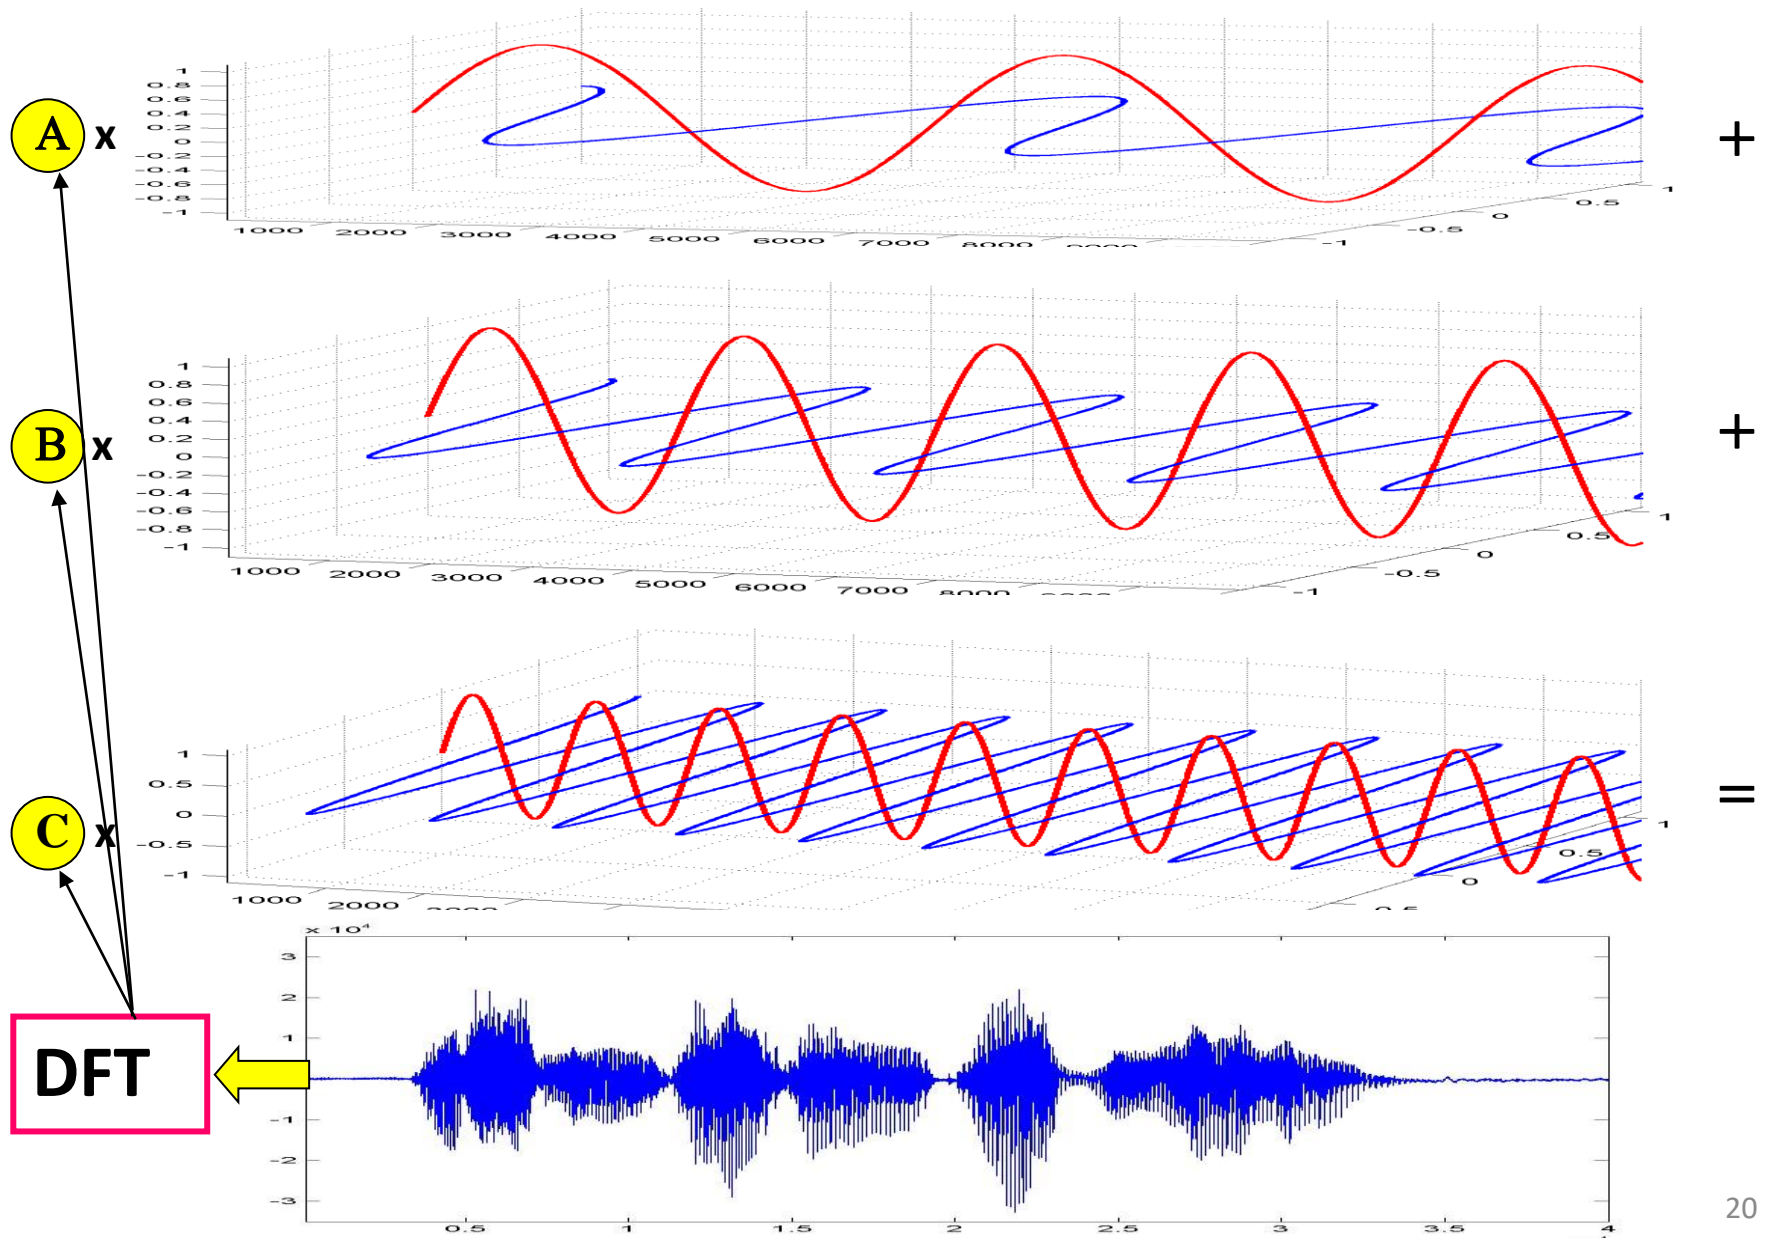
- Transform analysis: Decompose a sequence of numbers into a weighted sum of other time series
  - The component time series must be defined
  - For the Fourier Transform, these are complex exponentials
- The transform analysis determines the weights of the component time series



# The discrete Fourier transform




# The discrete Fourier transform



# A clean speech signal and it's spectrogram

Tom Sullivan spells out his last name: S U L L I V A N

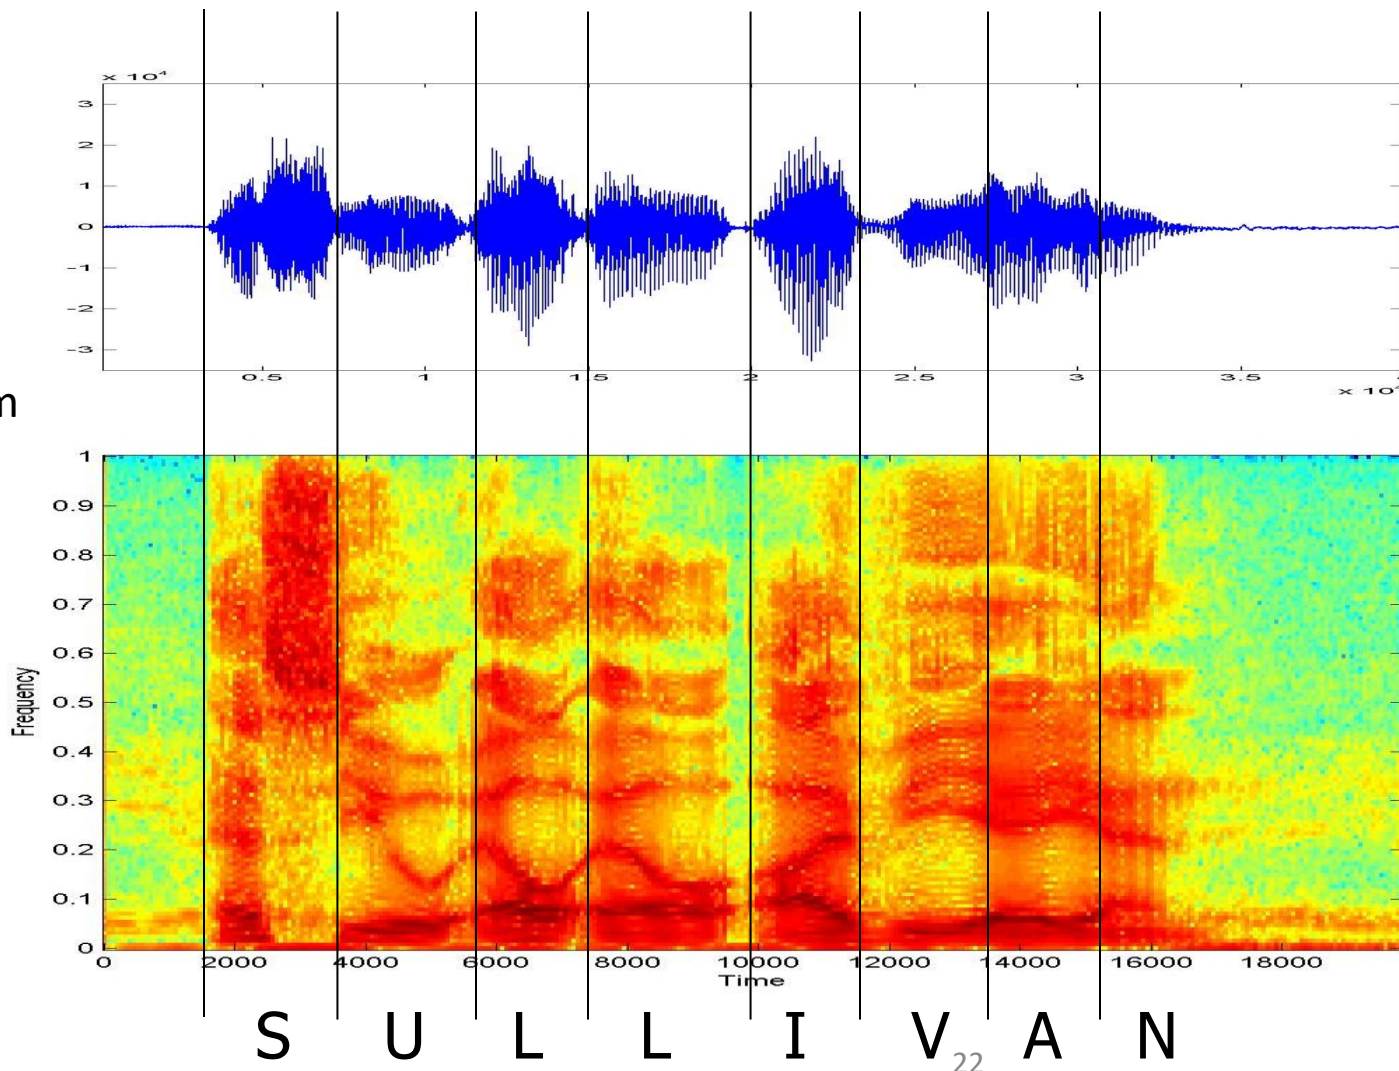






# A clean speech signal and it's spectrogram

Tom Sullivan spells out his last name: S U L L I V A N

**Tool:** Many free tools available to compute and study spectrograms, e.g. Audacity Praat





# Spectrum



- Audio/speech spectrum

<http://www.bnoack.com/index.html?http&&>  
[www.bnoack.com/audio/speech-level.html](http://www.bnoack.com/audio/speech-level.html)

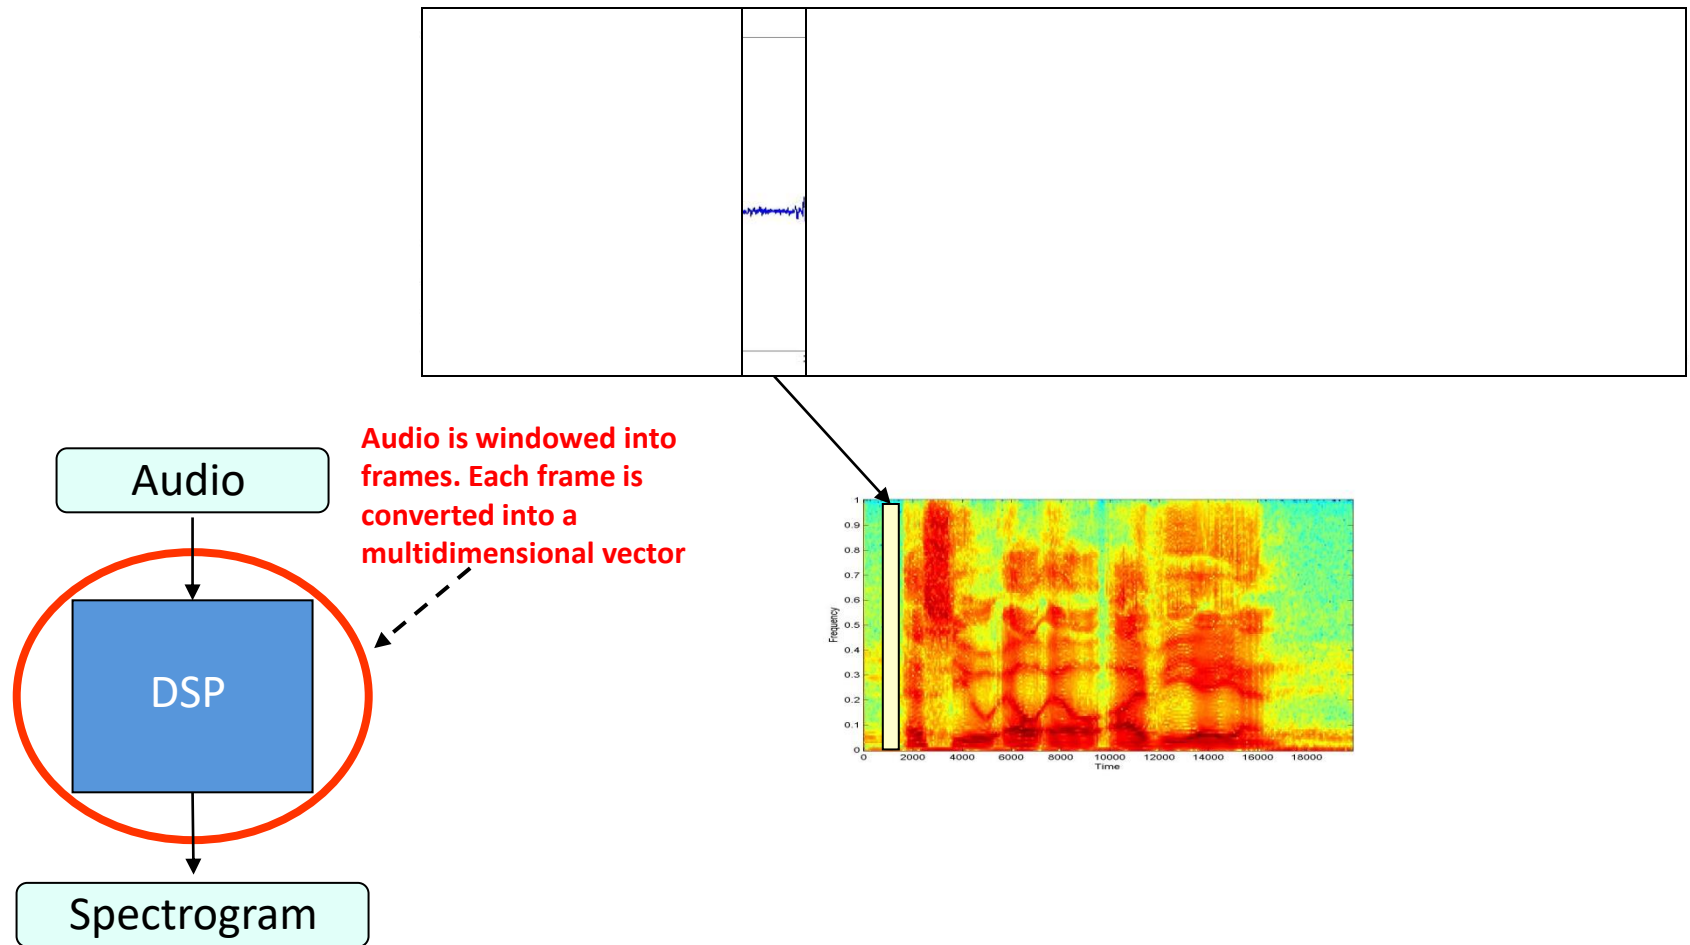
# Audio (sound) signatures in a spectrogram



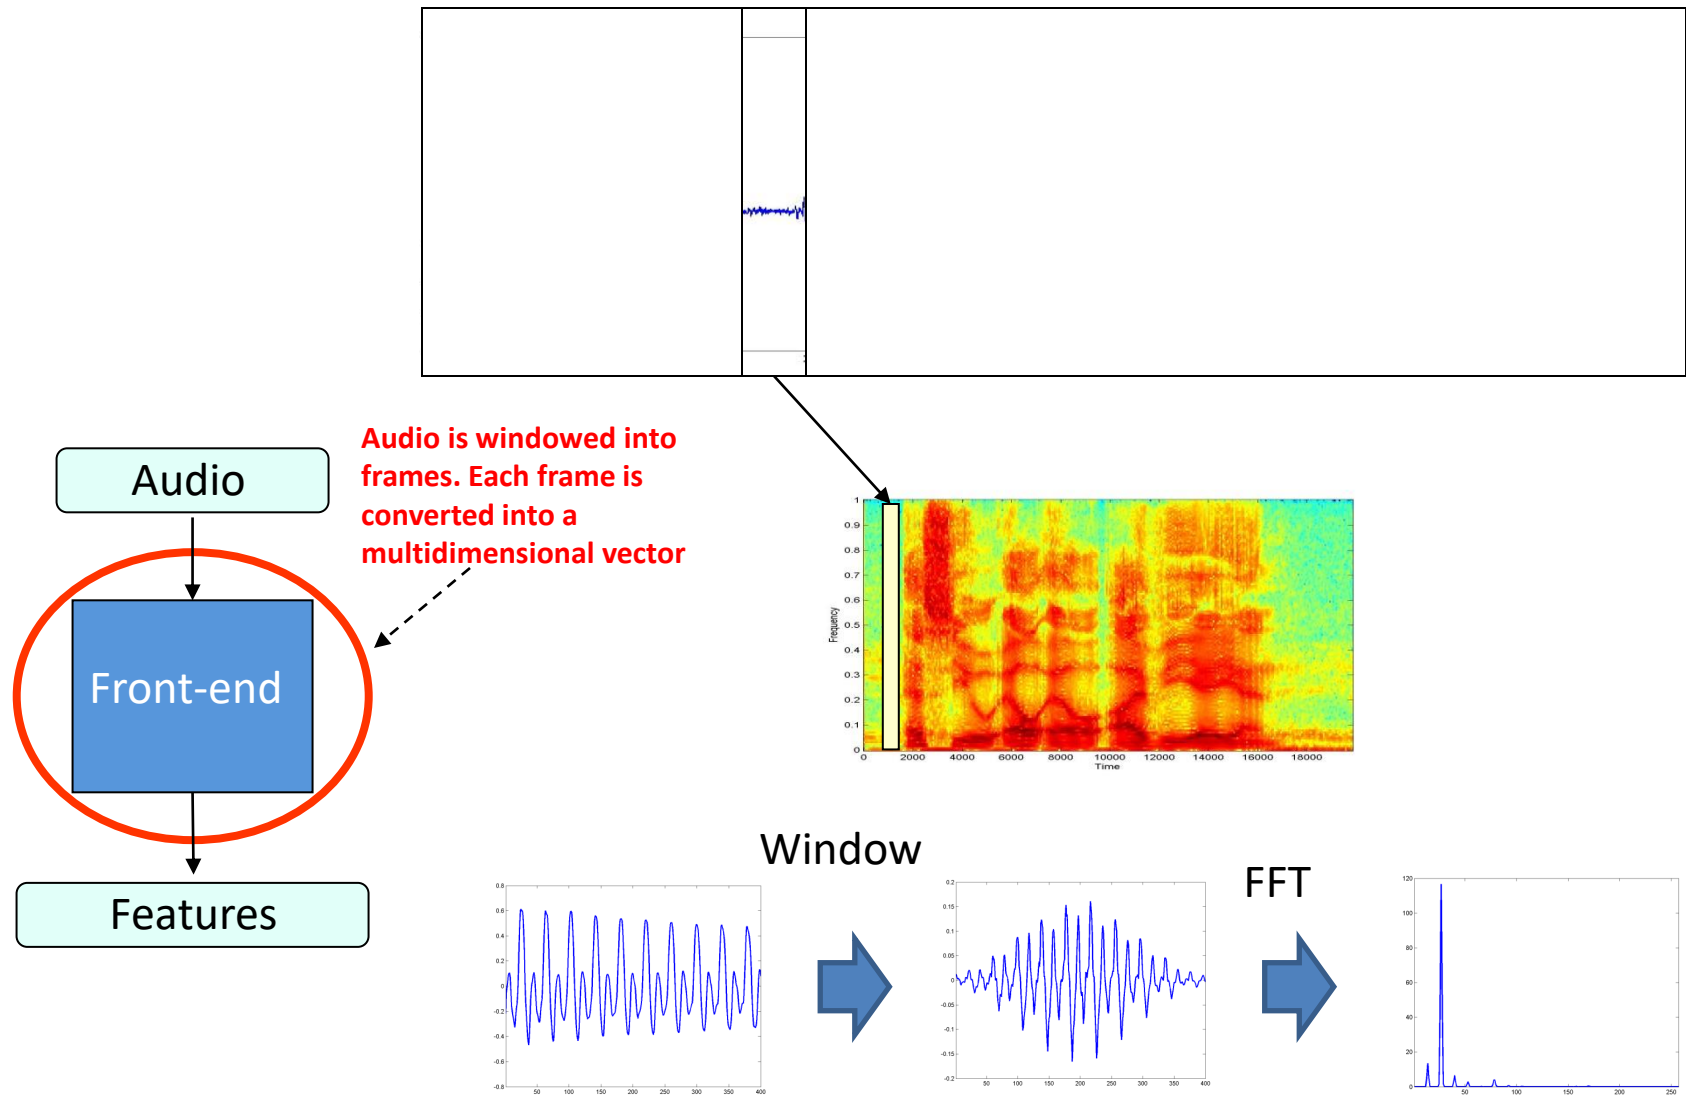
- A Ferrari going at 290kmph
- A barn owl



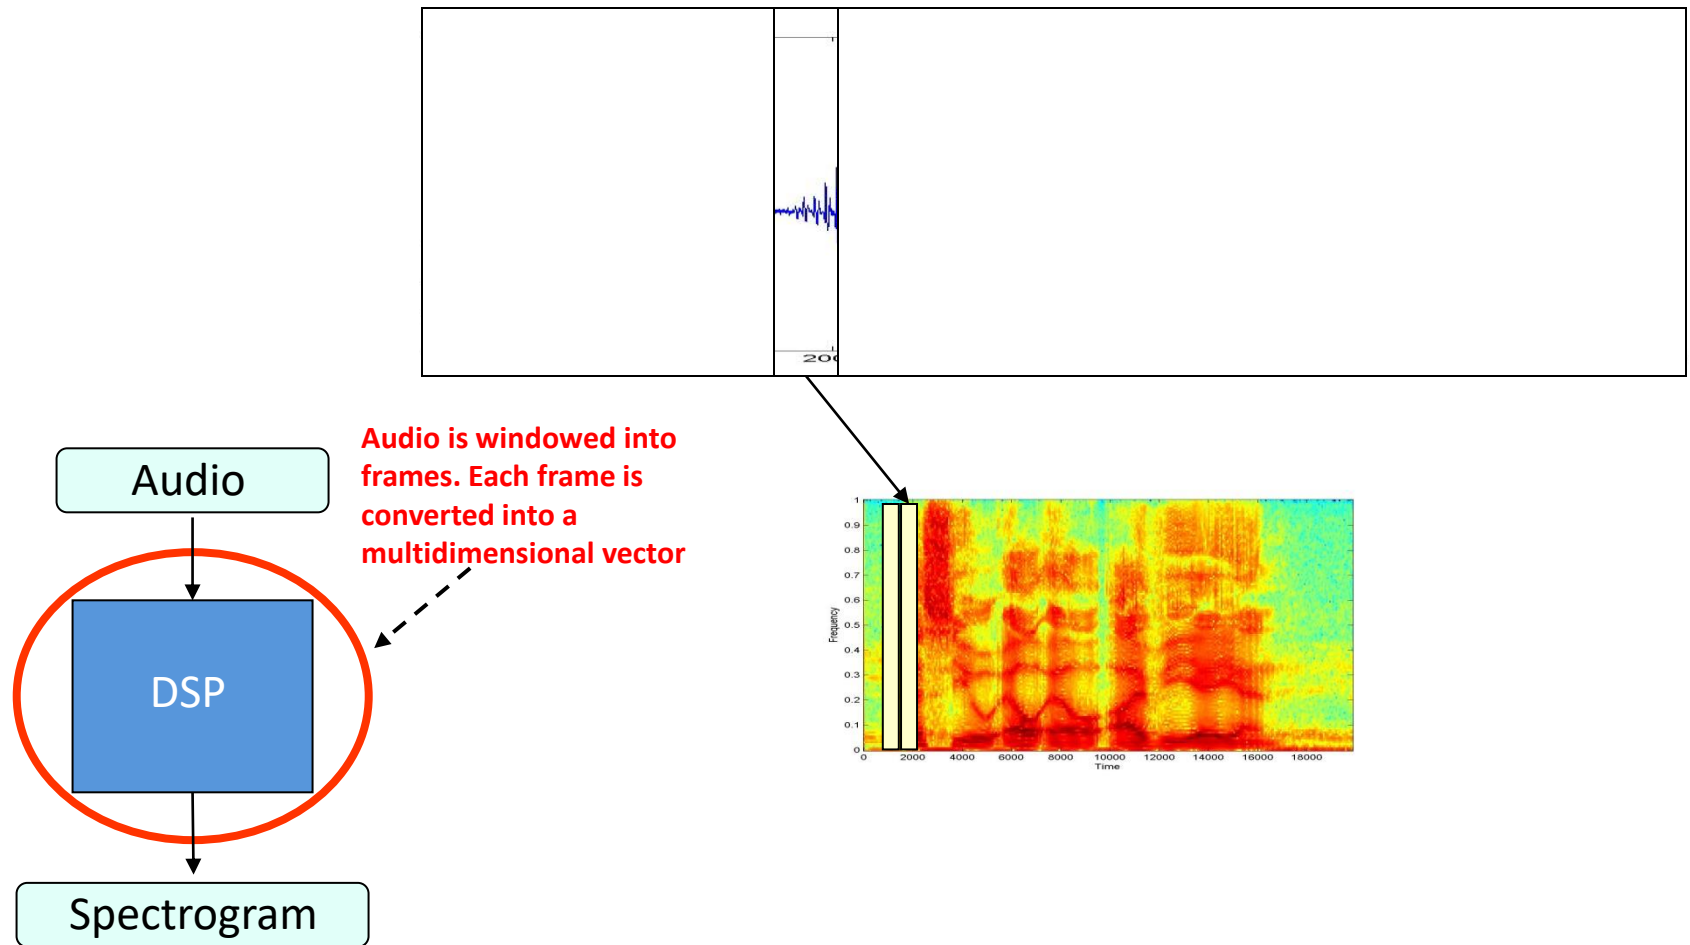
# How is the spectrogram computed?



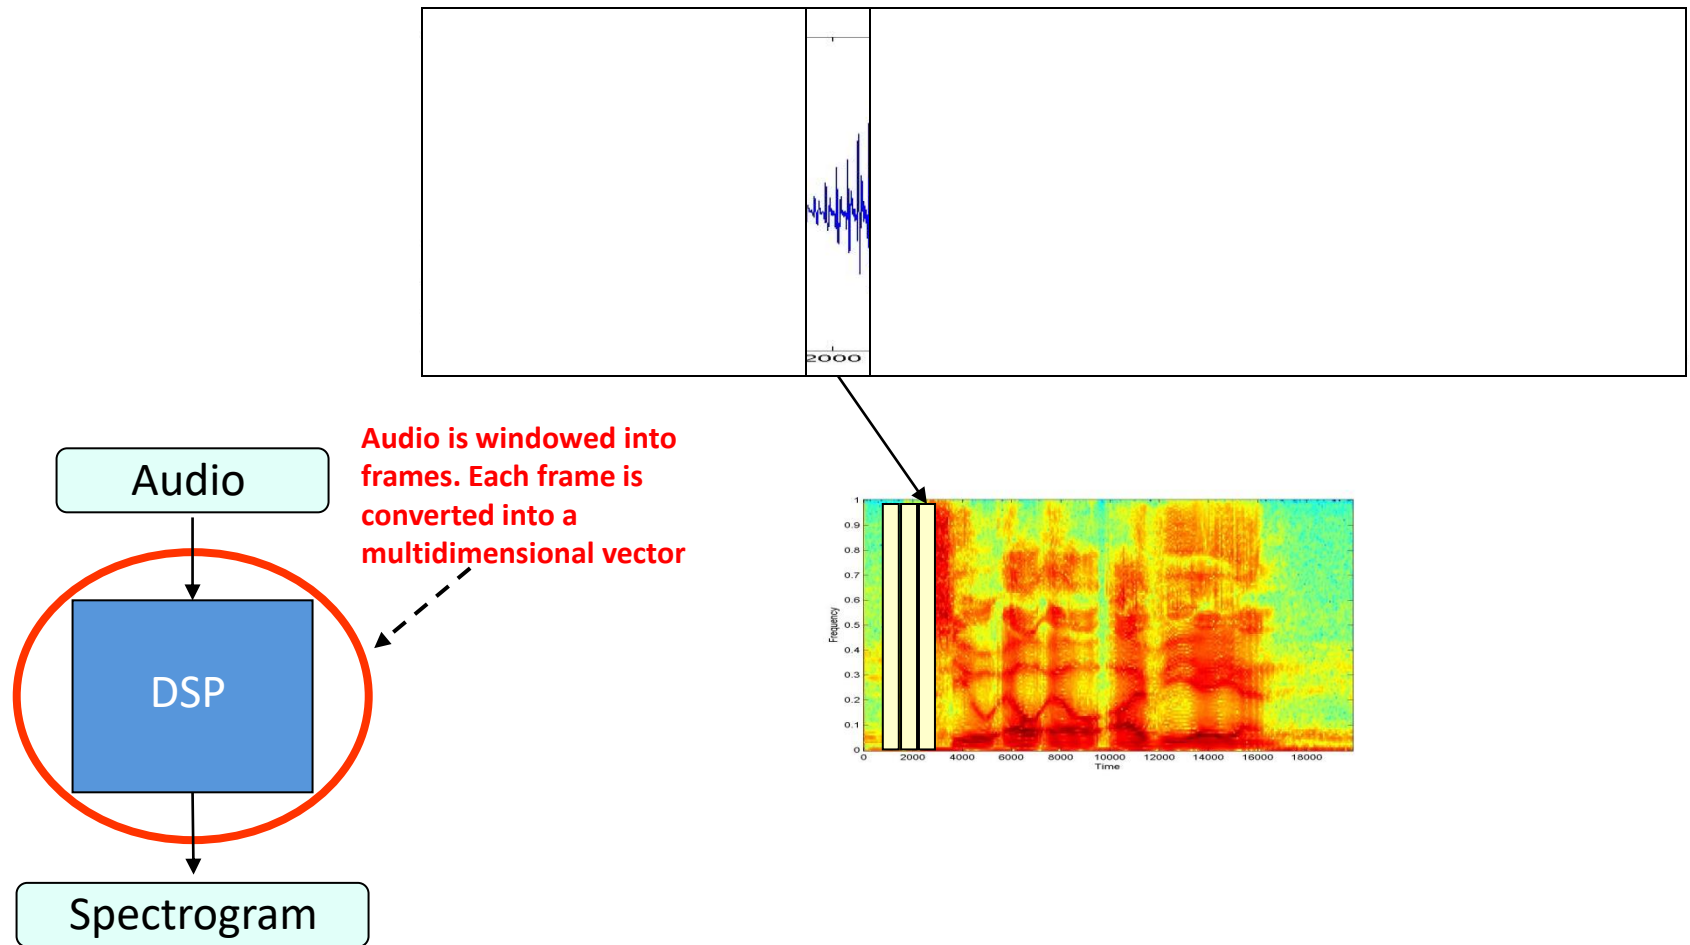
# How is this *vector* computed?



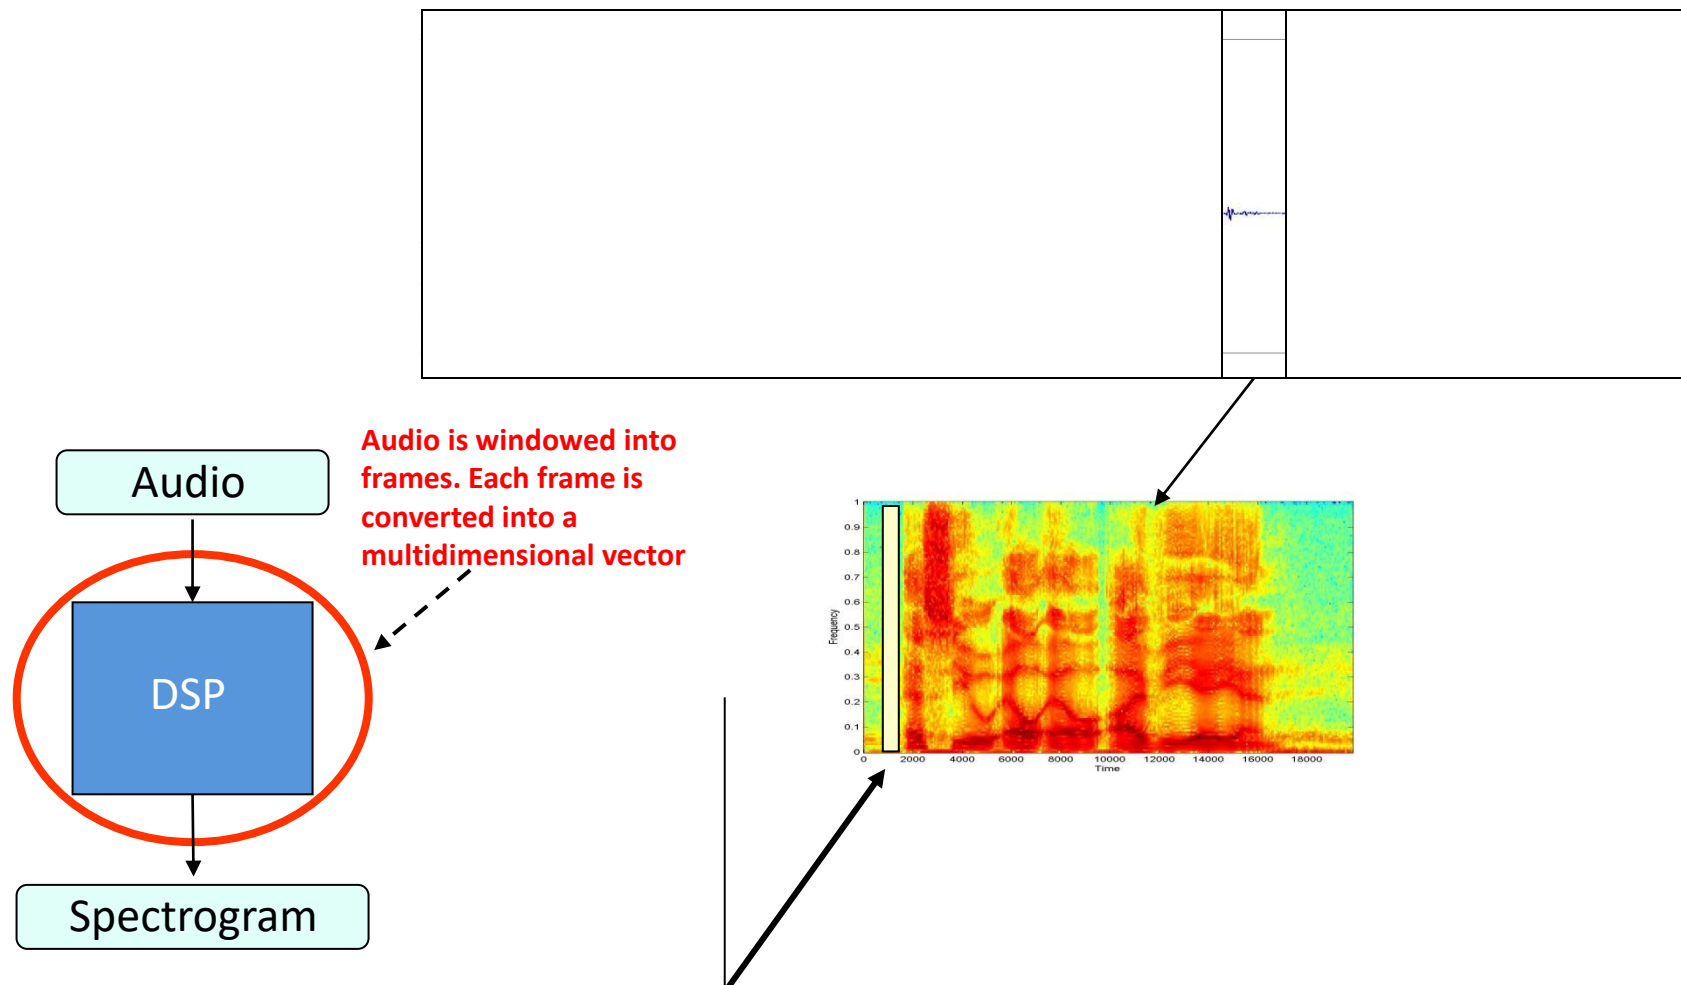
# How is the spectrogram computed?



# How is the spectrogram computed?




# A spectrogram is composed of vectors



- The “picture” is simply a visualization of the sequence of vectors
- Each high-dimensional vector represents a slice of time
- Each component of the vector represents a particular frequency at that time


# Computing Mel Cepstra



**The signal is processed in segments.  
Segments are typically 25 ms wide.**

**Adjacent segments typically overlap  
by 15 ms.**



# Computing Mel Cepstra



**The signal is processed in segments.  
Segments are typically 25 ms wide.**

**Adjacent segments typically overlap  
by 15 ms.**

# Computing Mel Cepstra




Segments shift every 10 milliseconds

Each segment is typically 20 or 25 milliseconds wide  
Speech signals do not change significantly within this short time interval



# Computing Mel Cepstra



Each segment is preemphasized


**Preemphasized segment**

The preemphasized segment is windowed


**Preemphasized and windowed segment**

# Pre-emphasizing a speech signal

- The spectrum of the speech signal naturally has lower energy at higher frequencies
- This can be observed as a downward trend on a plot of the logarithm of the magnitude spectrum of the signal
- For many applications this can be undesirable
  - E.g. Linear predictive modeling of the spectrum

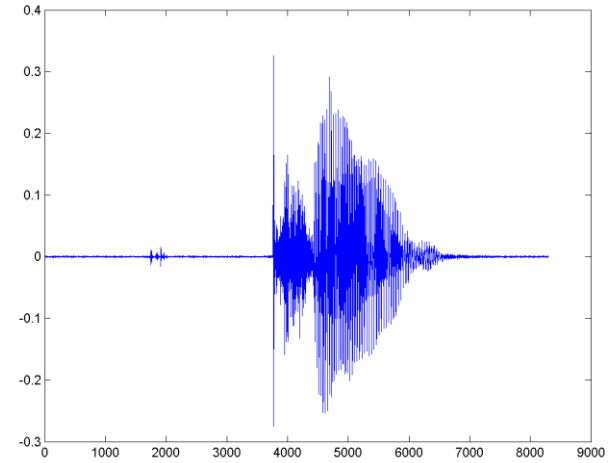


**Log(average(magnitude spectrum))**




# Pre-emphasizing a speech signal


- This spectral tilt can be corrected by preemphasizing the signal
  - $s_{\text{preemp}}[n] = s[n] - \alpha * s[n-1]$
  - Typical value of  $\alpha = 0.95$
- This is a form of differentiation that boosts high frequencies
- This spectrum of the preemphasized signal has more horizontal trend
  - Good for linear prediction and other similar methods




**Log(average(magnitude spectrum))**




# Windowing



- The DFT of *any* sequence computes the Fourier series for an infinite repetition of that sequence
- The DFT of a partial segment of a sinusoid computes the Fourier series of an infinite repetition of that segment, and not of the entire sinusoid
- This will not give us the DFT of the sinusoid itself!




# Windowing






- In addition, the implicit repetition of the observed signal introduces large discontinuities at the points of repetition
  - This distorts even our measurement of what happens at the boundaries of what has been reliably observed
  - The actual signal (whatever it is) is unlikely to have such discontinuities

# Windowing




- While we can never know what the signal looks like outside the window, we can try to minimize the discontinuities at the boundaries
- We do this by multiplying the signal with a *window* function
  - We call this procedure windowing
  - We refer to the resulting signal as a “windowed” signal

# Windowing




# Windowing



## Geometric windows:

- Rectangular (boxcar)
- Triangular (Bartlett)
- Trapezoid



## Cosine windows:

- Window length is  $M$
- Index begins at 0


Hamming:  $w[n] = 0.54 - 0.46 \cos(2\pi n/M)$

Hanning:  $w[n] = 0.5 - 0.5 \cos(2\pi n/M)$

Blackman:  $0.42 - 0.5 \cos(2\pi n/M) + 0.08 \cos(4\pi n/M)$




# Computing Mel Cepstra (MFCC)



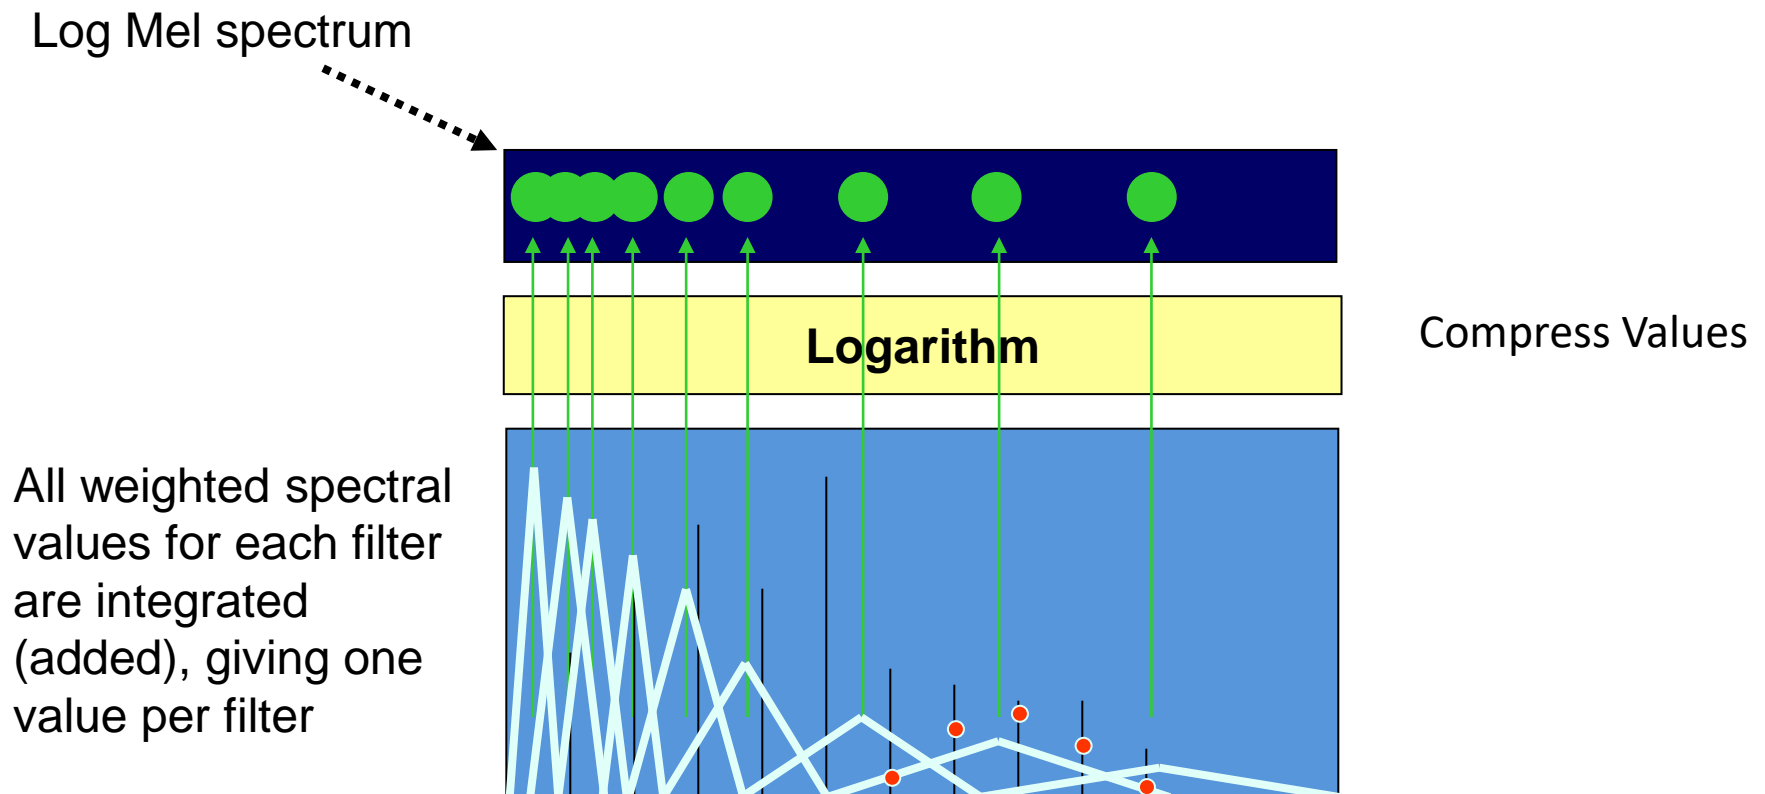
# The process of parametrization

All weighted spectral values for each filter are integrated (added), giving one value per filter

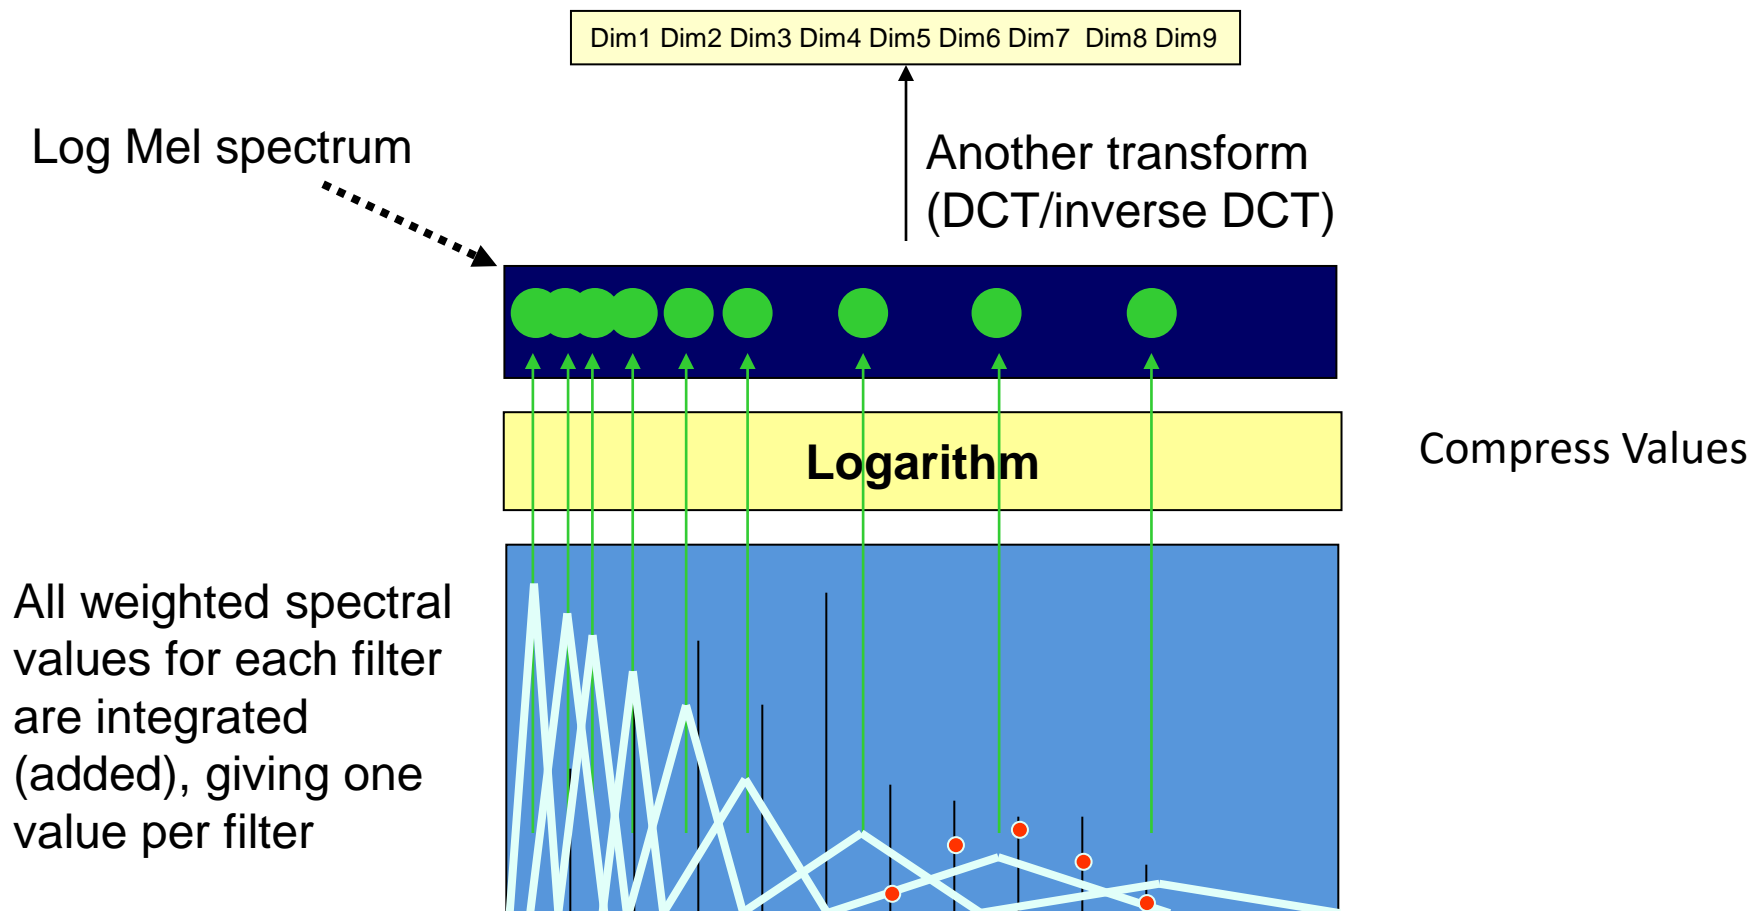


Compress Values

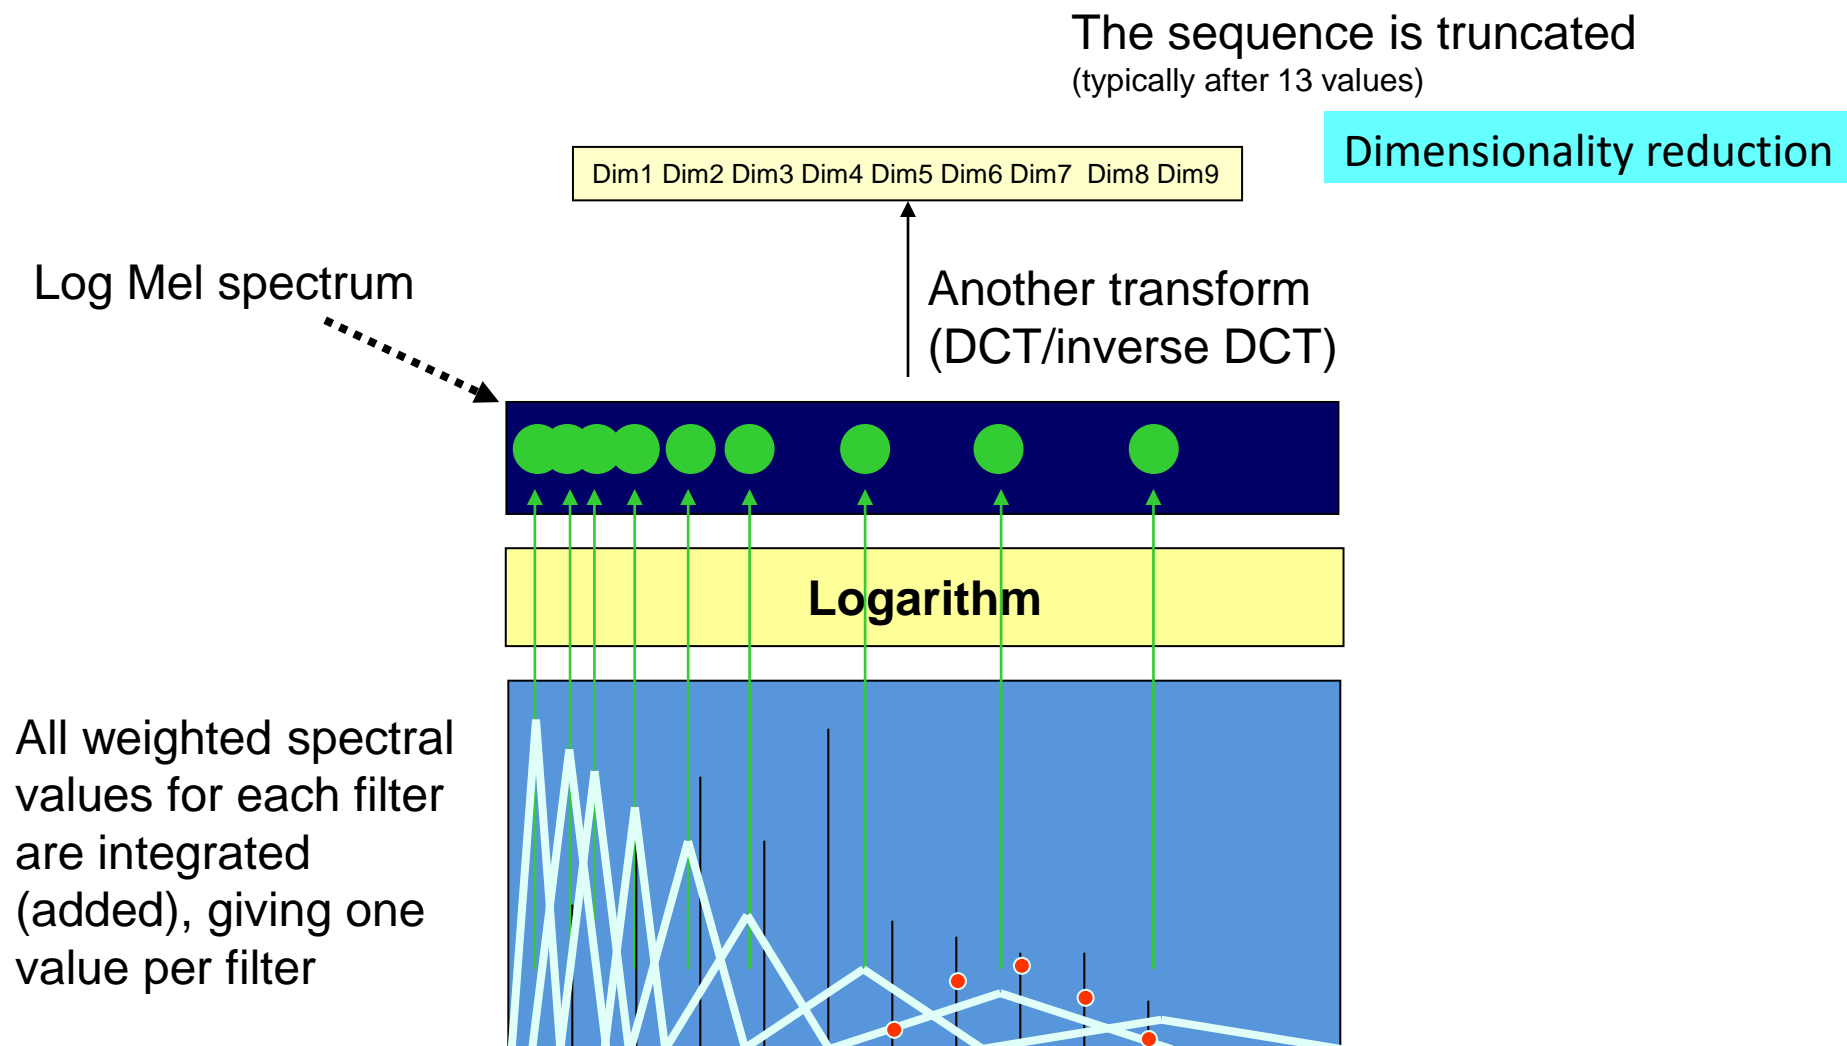
# The process of parametrization



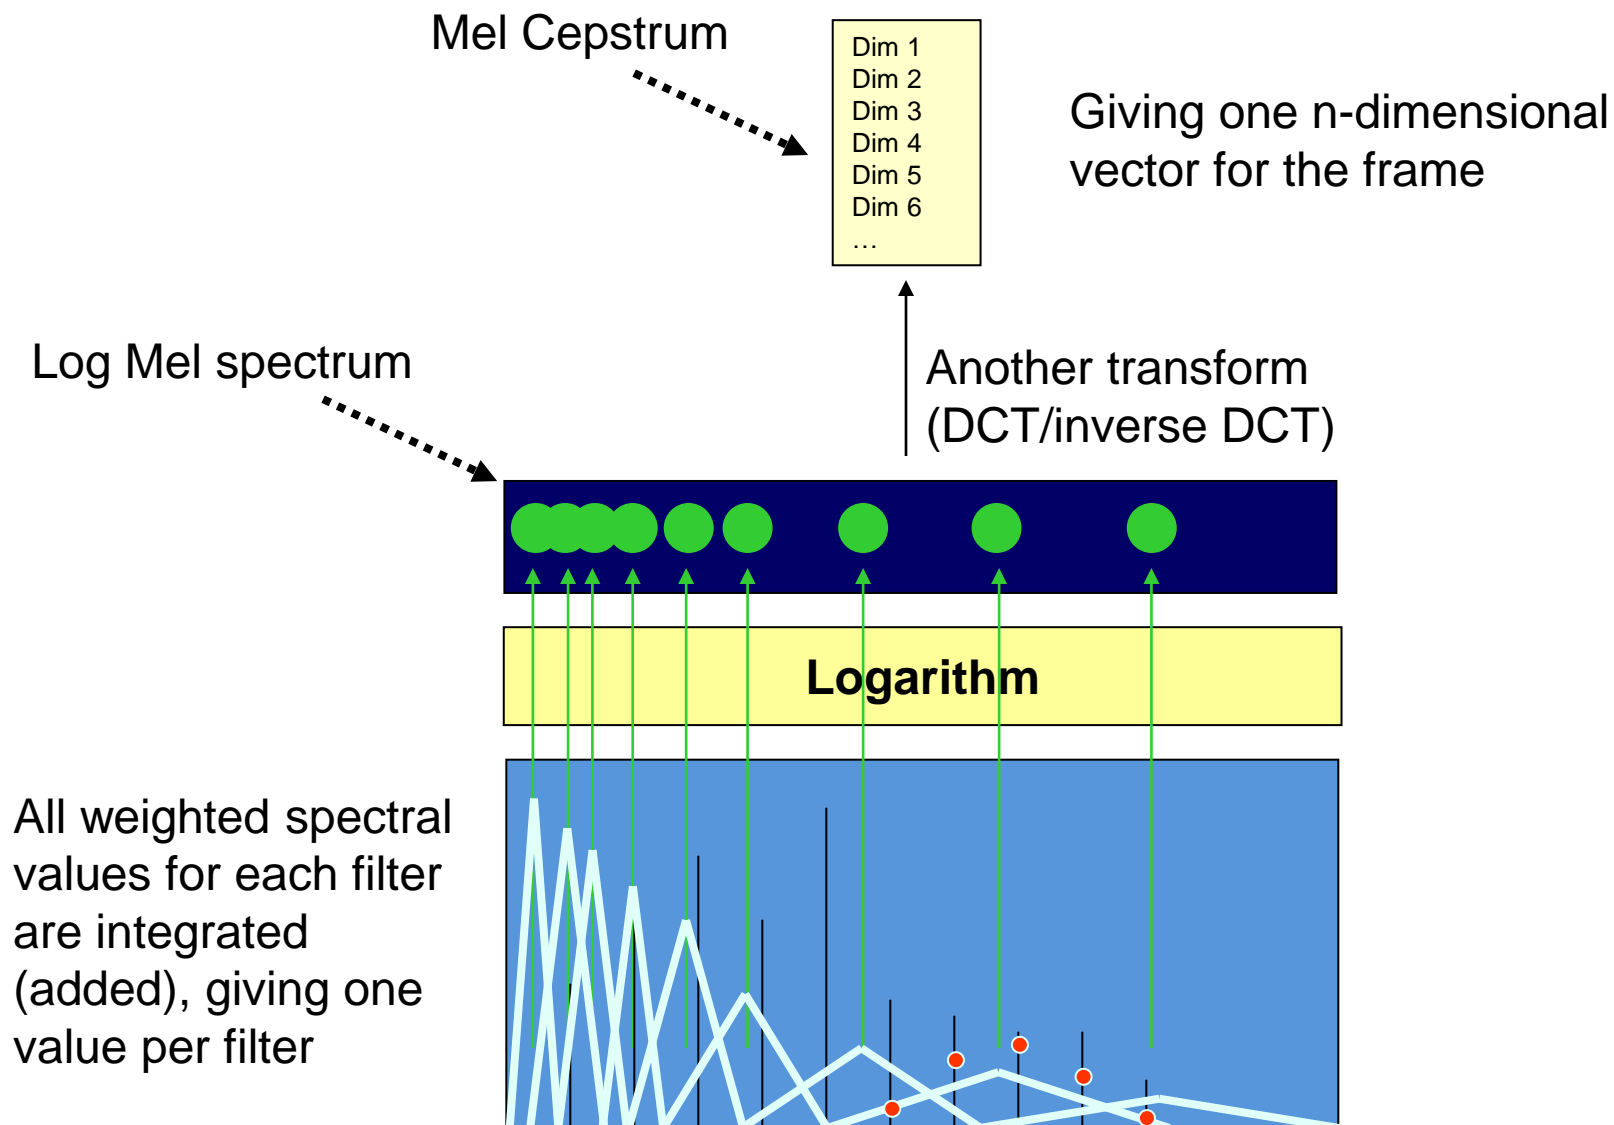
# The process of parametrization



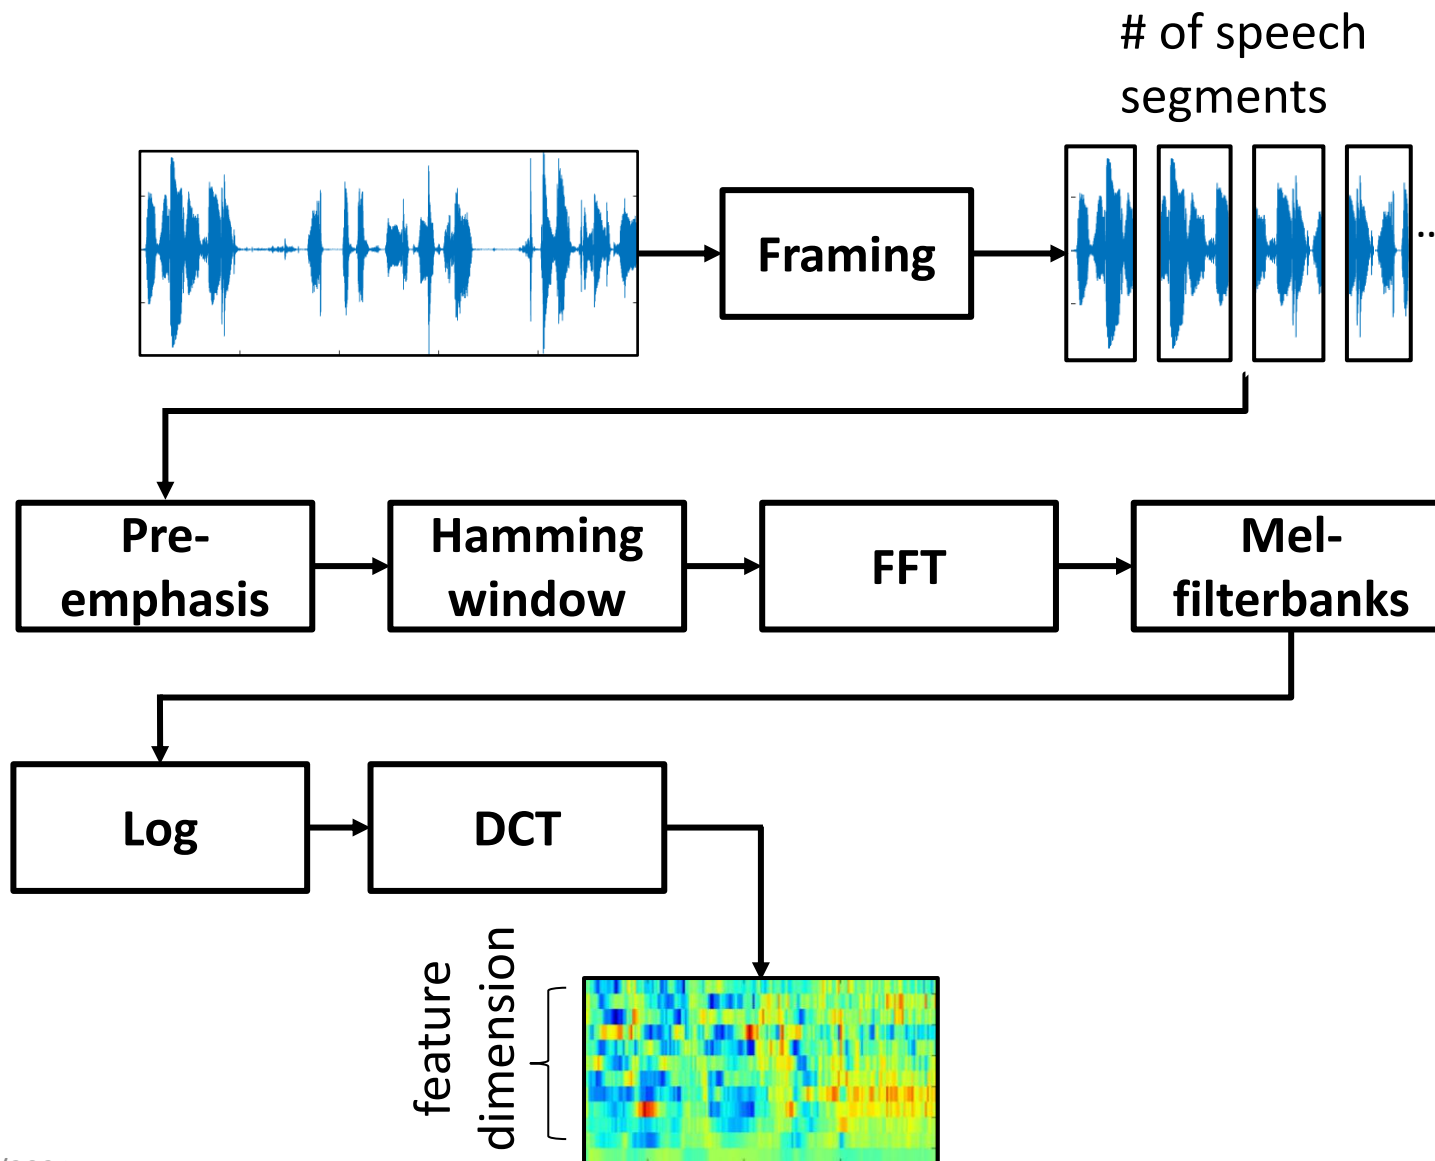
# The process of parametrization










# The process of parametrization



# Extracting cepstra





# An example segment






# The process of parametrization



**The entire speech signal is thus converted into a sequence of vectors. These are Mel cepstral vectors (also known as MFCC, or Mel Frequency Cepstral Coefficients).**

# In general, audio is represented as features



## Step 1

- The time domain audio signal is transformed into a sequence of different measurements
  - These are usually vectors, and are called feature-vectors
  - The process is called **parametrization**

## Step 2

- Either the features are used directly or
- Secondary features are extracted from them
  - Secondary features are then compared

# Variations to the basic theme

- Perceptual Linear Prediction (PLP) features:
  - ERB filters instead of MEL filters
  - Cube-root compression instead of Log
  - Linear-prediction spectrum instead of Fourier Spectrum
- Auditory features
  - Detailed and painful models of various components of the human ear


# Audio analysis overview

- Feature computation (selected)
  - Spectra
  - Spectrograms
  - Mel-cepstra
  - i-vectors
  - Supervectors
  - Bag-of-words
  - NN-based features
  - Visual features
  - Videographic features
- Specific applications of audio processing (selected)
  - Audio authentication
  - Audio enhancement
  - Audio fingerprinting
  - Audio localization
  - Audio object detection
  - Audio retrieval
  - Audio summarization
  - Environmental profiling
  - Geolocation
  - Source Identification
  - Source separation
  - Speaker identification
  - Speaker profiling
  - Speaker verification
  - Speech recognition
  - Speech separation
- Key analysis techniques (selected)

# Macro and micro features

- **macro-features**
  - derived from large windows, or the entire signal
  - Derived as secondary features by aggregating features derived from short-duration windows
    - typically 25 ms wide
    - Adjacent segments typically overlap by 15 ms
- **micro-features**
  - derived from very small (duration  $< 25\text{ms}$ ) segments of the signal

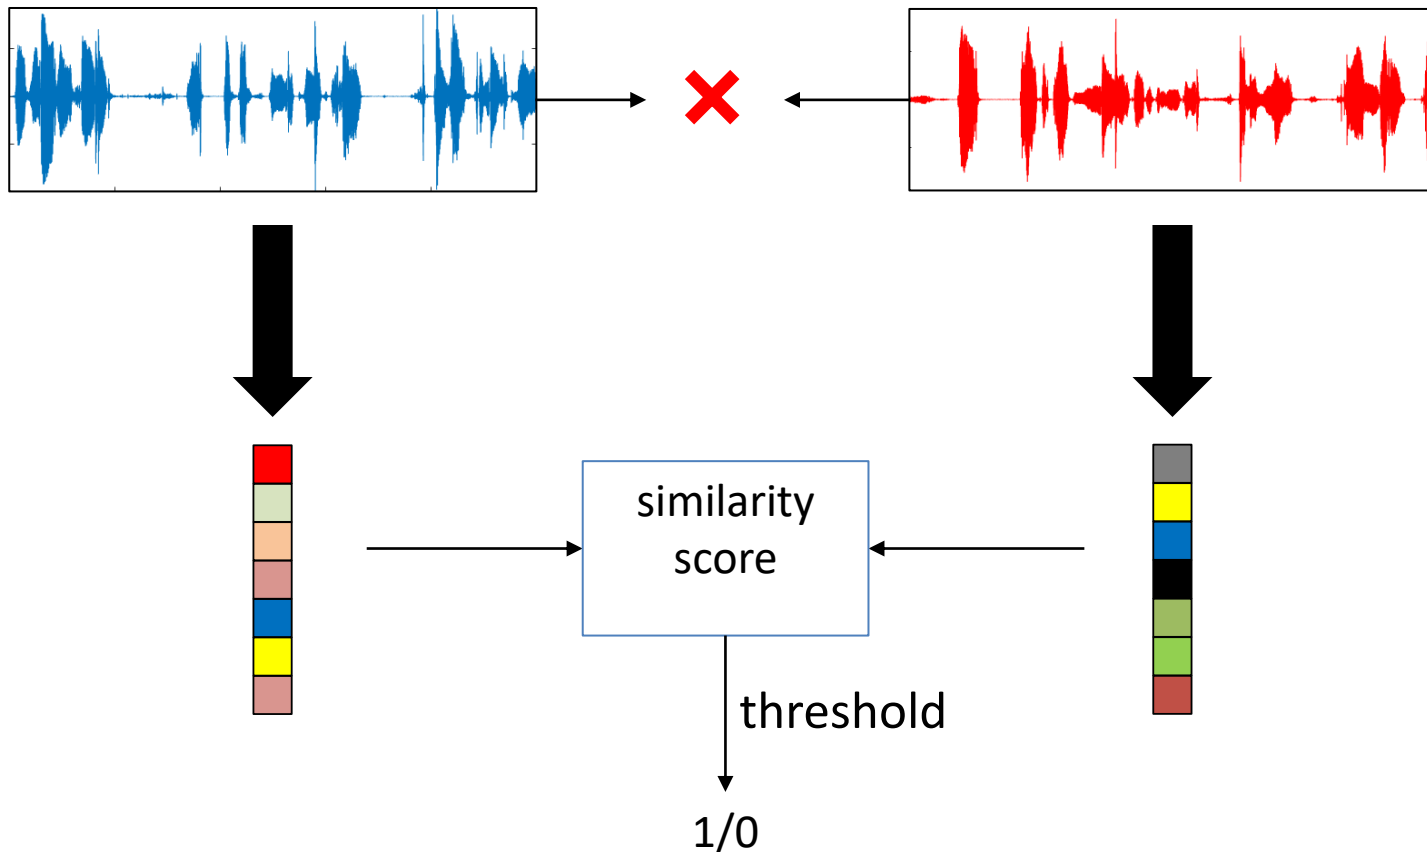
# Audio/Speaker Matching




- Determine if the speaker in a “test” recording is the same as that in a previously obtained recording

# Approach

- Extract a “representation vector” from the recordings
- Identify matches by comparing the representation vectors
- **Key challenge: Learning the right representations for the recordings**




# Extracting a representation vector from a recording



- Must focus on *overall* characteristics of the recording
  - Rather than instantaneous patterns, which may never be repeated
- Must reduce *variable-sized* recordings to a fixed-size vector
  - Required for comparisons




# Features for matching




- Convert the audio recording to a sequence of *spectral* vectors, that together comprise a “spectrogram”
  - May applications use *cepstral* vectors instead, which are derived from the spectral vectors
- Estimate the *distribution* of the vectors in the recording
  - **Hypothesis:** The ID of the speaker is captured by this distribution
    - Modes capture both linguistic and acoustic tendencies

# Features for matching



- Learn a Gaussian mixture density from the collection of *cepstral* vectors in the recording
  - Adapt a “universal background model” to do so
- Concatenate the parameters of the GMM into a single vector called a “**Super-Vector**”
- Reduce the dimensionality of the super-vector through discriminative factor analysis
  - “**Probabilistic Latent Discriminant Analysis**” (PLDA)

# Supervectors and I-vectors




- In practice, there is generally insufficient data in the recordings to compute a distribution
  - So we adapt a “Universal” model to the individual recordings
- The adapted distribution captures the overall statistical characteristics of the recording
- We can represent this using a *single* vector, obtained by simply concatenating the means of the individual Gaussians in the mixture

$$S = [\mu_1 \quad \mu_2 \quad \dots \quad \mu_L]$$

- $\mu_i$  is the mean of the  $i$ th Gaussian
  - $L$  is the total number of Gaussians
- This vector is called a **Supervector** representing the distribution of spectral vectors derived from the recording


# Statistical Distribution Matching



## Problems with this approach

- Needs enough data to estimate distribution
  - Even with all the dimensionality reduction, needs several minutes in each recording for reliable results
- Average's information across time: Ignores temporal structure.
  - More generally, ignores structural/phonetic nature of speech,
- No reliable mechanism to *incrementally* include evidence from additional data
- Not robust to changes due to age, disease etc.

# Statistical Distribution Matching




- To further reduce data requirement, various dimensionality-reducing techniques like **factor analysis** and **Linear Discriminant Analysis (LDA) (for classification)** are employed
- With jargon such as **I-vectors**, "**Total variability space**", **PLDA**..

## ***l-vectors: concept***


- “l”-vectors are derived from super-vectors through “factor analysis”
- Factor Analysis is in turn an extension of “Principal Component Analysis”..

# Principal Component Analysis



- In high dimensions, most data lie on or close to lower dimensional “linear manifold” -- a hyperplane
- Principal Component Analysis attempts to find this hyper plane and “place” the data on it


# Principal Component Analysis



- PCA finds a plane such that the total squared distance from the points to their “shadow” on the plane is minimum
  - The “shadow” is the “projection” of the data on the plane
- The “shadow” is assumed to be the true data, the off-plane component is the noise



# Principal Component Analysis




- PCA model:

- The actual data are formed on the plane (red arrow on plane)
- The noise is *orthogonal* to the plane

$$Y = BX + N$$

- $Y$  is final vector (3D vector in our example, indicated by red ball)
- $B$  is a matrix with the unit vectors for the “bases” (long grey arrows, coordinate axes of plane,
  - 3x2 matrix in our example: 2 vectors in 3 dimensions
- $X$  represents coordinates of the shadow point (2D vector in our example)
- $N$  is the noise (blue line)

# Principal Component Analysis




- **PCA statistical generative model:**

$$Y = BX + N$$


- $X$  is drawn from a 0-mean Gaussian distribution with a diagonal covariance matrix
  - Co-ordinates are uncorrelated
- $N$  is drawn from a 0-mean Gaussian with a **low-rank** covariance
  - Rank is 1 in our example
  - More generally, for D-dimensional data, explained through K-dim PCA, rank of noise covariance is D-K
- **PCA challenge:** *Given a collection of  $Y$  vectors find bases  $B$  of the data plane and the coordinates  $X$  on the plane for each vector  $Y$* 
  - Find the grey lines, and the location of the shadow for each point
- Obtained through a maximum-likelihood estimator, which is your familiar PCA

# Inadequacy of PCA




- Assumes noise is always *orthogonal* to the data
- Not a reasonable assumption in most cases
  - e.g. in speech, noise may also sound a bit like speech
  - i.e. noise is not perpendicular to the plane, although it does not *lie* on the plane

# Factor Analysis



- Factor analysis:
  - The noise is not required to be perpendicular to the plane
    - May be at any angle to it
- Factor analysis model for the data:
  - Data are formed on the plane
  - A random noise is added to the data

# Factor Analysis




- Factor Analysis statistical generative model:

$$Y = BX + N$$

- $X$  is drawn from a 0-mean Gaussian distribution with a diagonal covariance matrix
- $N$  is drawn from a 0-mean Gaussian with a **full-rank** covariance
  - Rank is 3 in our example
  - More generally, for D-dimensional data, explained through K-dim FA, rank of noise covariance is D
- **FA challenge:** *Given a collection of  $Y$  vectors find bases  $B$  and the coordinates  $X$  on the plane for each vector  $Y$*
- The coordinates  $X$  for any  $Y$  are called the *factors* of  $Y$
- Must be estimated through an iterative *expectation maximization algorithm*

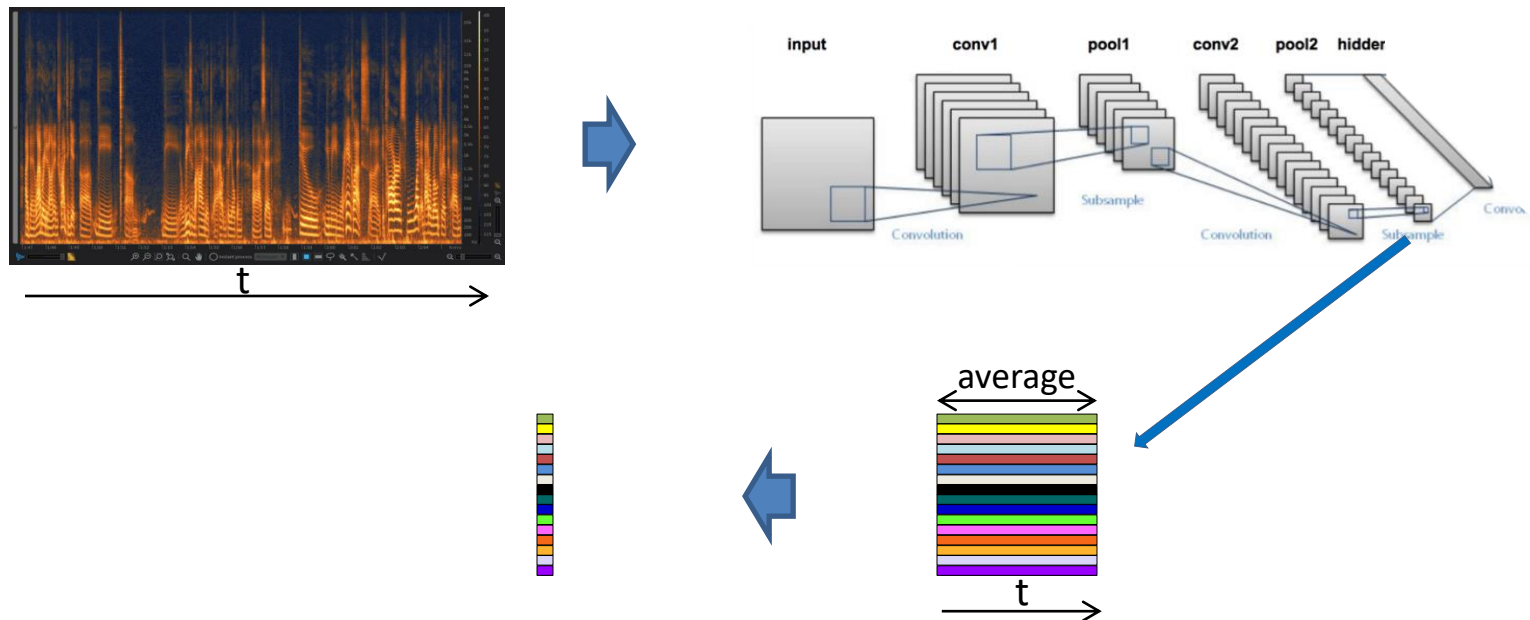
# I-vectors applied to speaker identification



$$S = [\mu_1 \quad \mu_2 \quad \dots \quad \mu_L]$$

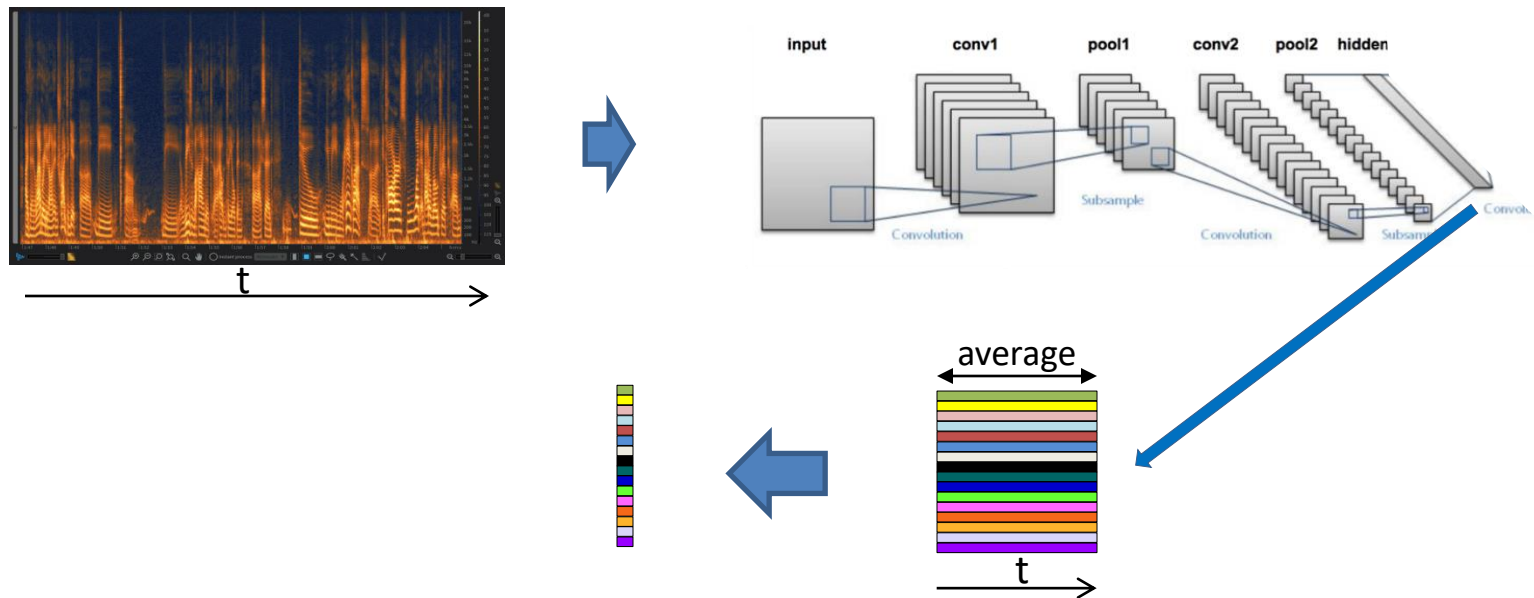
- Each recording from every speaker gives us a single super-vector
- Factor analysis is applied to the collection of super-vectors
- The **factor** vector corresponding to each super-vector is called the “I”-vector for that recording
- I-vectors of two recordings can be directly matched to determine if they are from the same speaker or not
  - Matching is done using machine learning algorithms

# CNNs for feature extraction



- Treat the spectrogram as an image
- Run a “convolutional neural network” over the image
  - Will give you several scaled-down versions of the image, called “maps” as output
- Average the final maps across time to derive a fixed-size feature


# Neural Networks/CNNs for feature extraction




- “Filters” capture local time-frequency patterns
  - Good! Capture structural characteristics
  - Bad: Only retains average occurrence frequency of patterns
- No mechanism for incremental inclusion of new information




# Problem from a mathematical perspective: how the CNN is trained



- The CNN is originally trained to classify between thousands of “training” speakers
  - Network has  $N$  outputs, where  $N$  = no. of training speakers
  - For any recording, only the output corresponding to that speaker must “light up”
- The final classification layer is then removed
  - The remaining network computes features that make it easy to classify between speakers

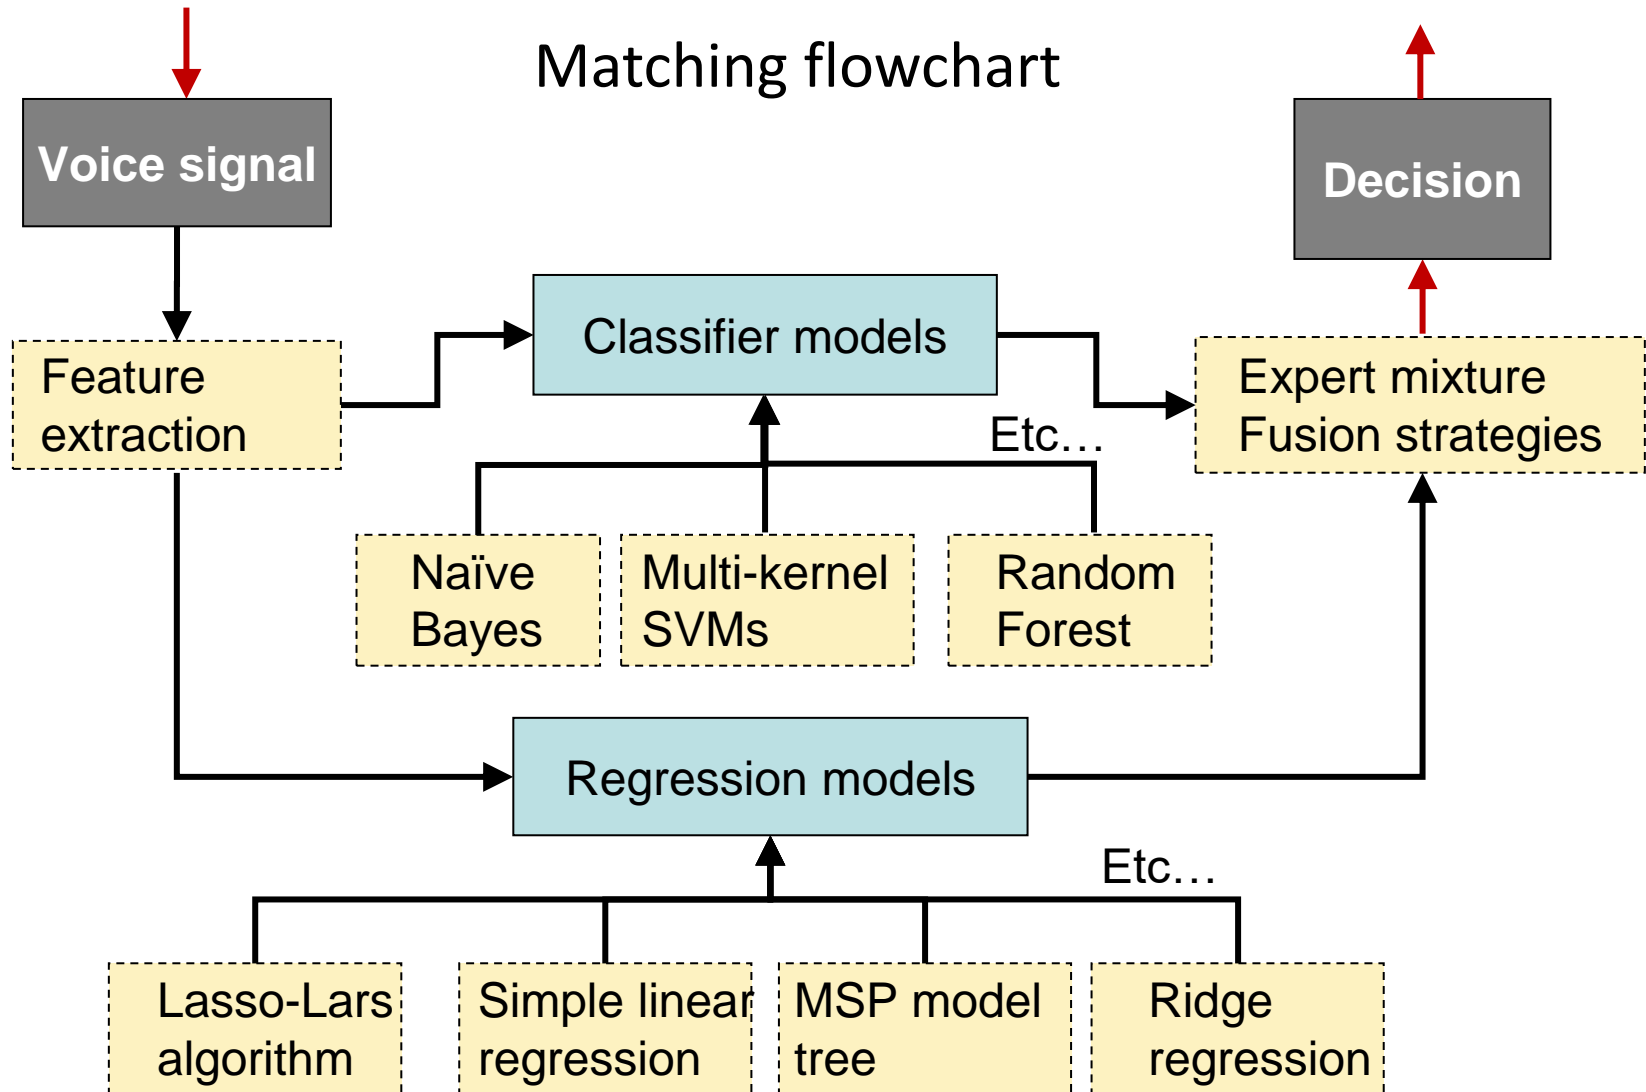


# Problem from a mathematical perspective

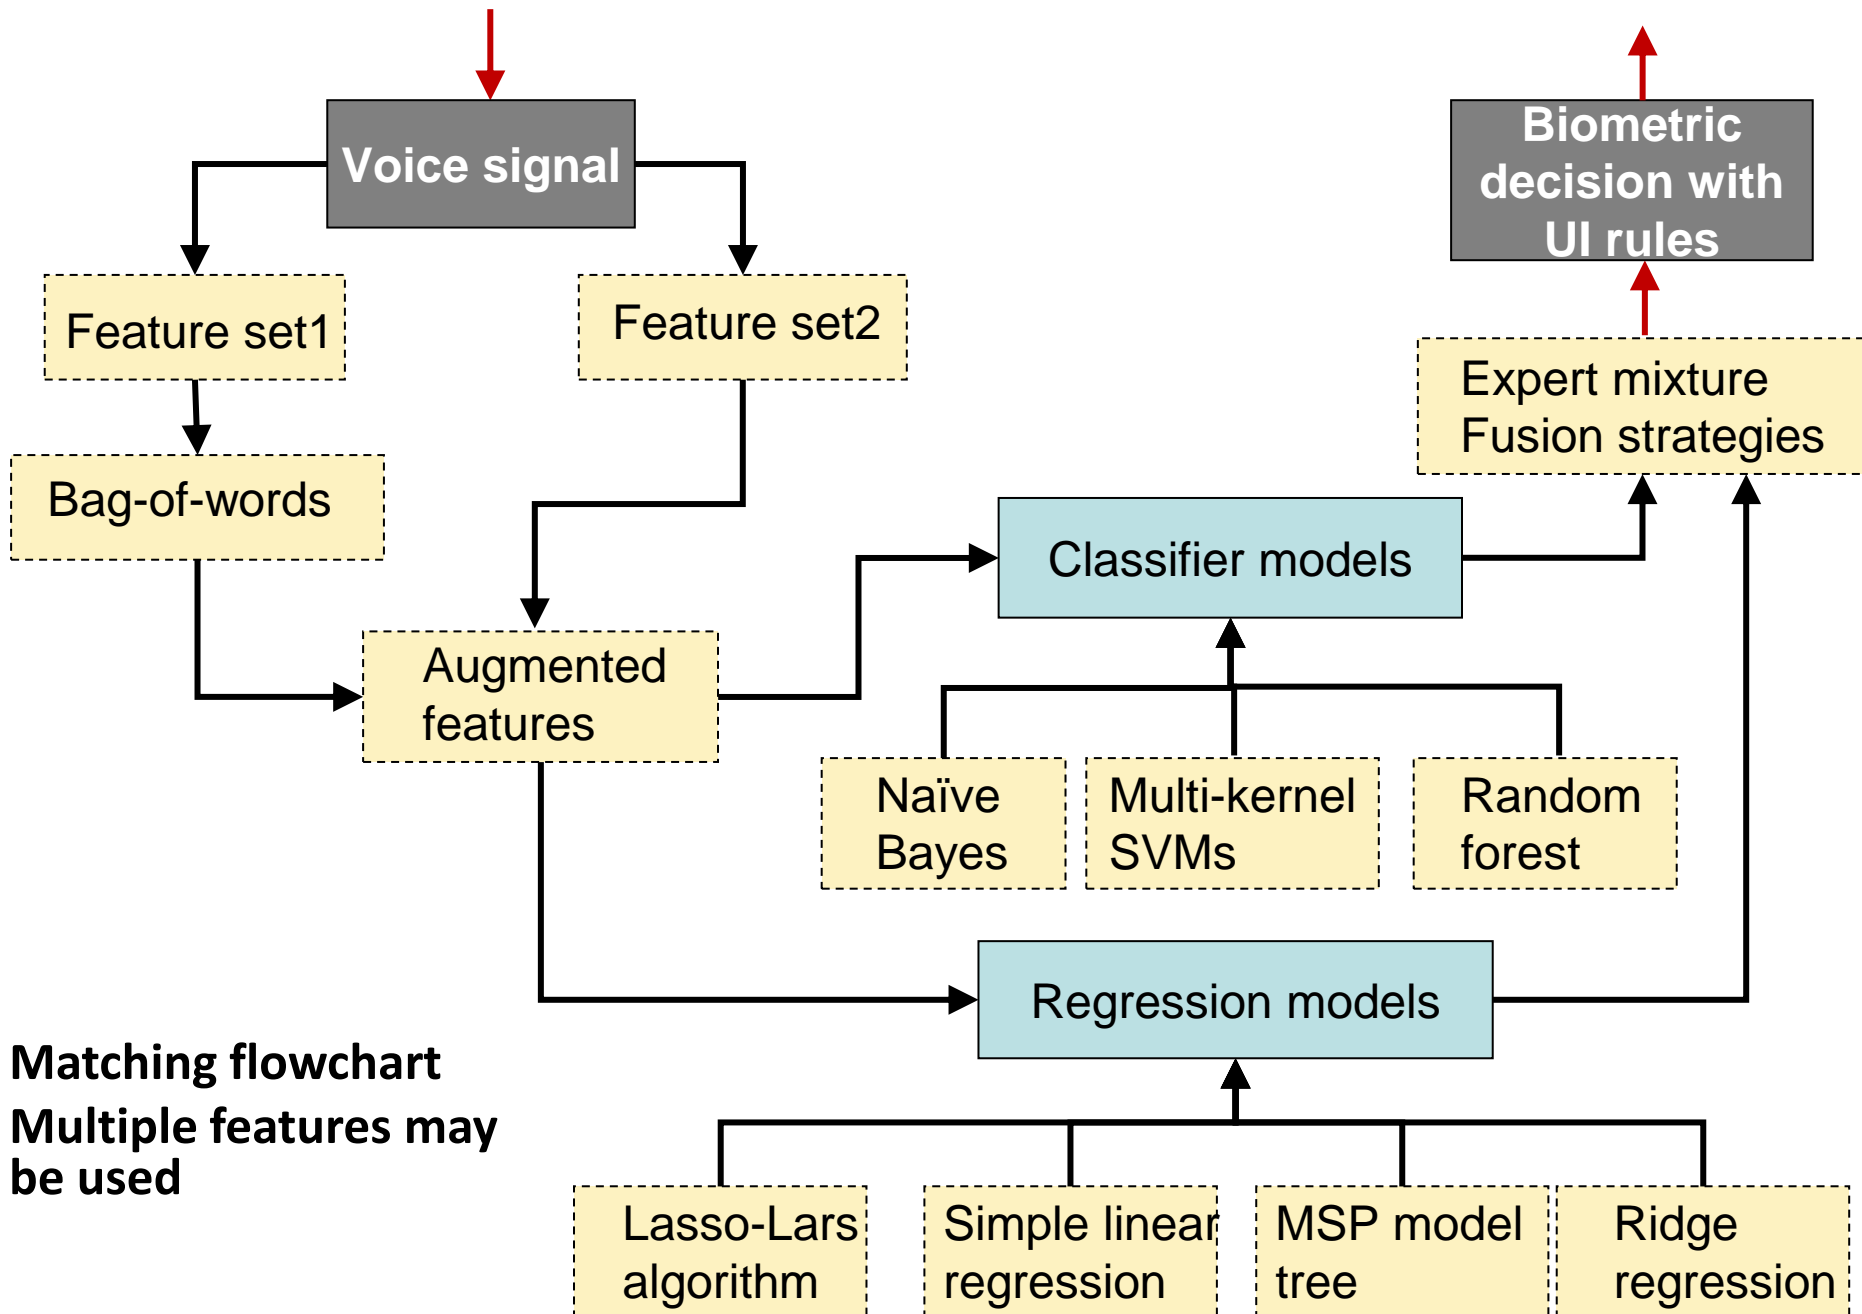


- A neural network computes class (speaker ID) log posterior:  $\log P(Y|x)$ 
  - Or alternately,  $P(Y|x)$
- Given a collection of **speech segments**  $x_1, x_2, \dots, x_T$  the sum output computes  $\sum_t \log P(Y|x_t)$ 
  - Or alternately,  $\prod_t P(Y|x_t)$
- What is *really* required
  - $P(Y|x_1, \dots, x_T)$
- But  $\prod_t P(Y|x_t) \neq P(Y|x_1, \dots, x_T)$
- The “average” operation is inappropriate to combine evidence from inputs

# From features to decisions



# From features to decisions



# So far...

- Feature computation (selected)
  - Spectra
  - Spectrograms
  - Mel-cepstra
  - i-vectors
  - Supervectors
  - Bag-of-words
  - NN-based features
  - Visual features
  - Videographic features
- Specific applications of audio processing (selected)
  - Audio authentication
  - Audio enhancement
  - Audio fingerprinting
  - Audio localization
  - Audio object detection
  - Audio retrieval
  - Audio summarization
  - Environmental profiling
  - Geolocation
  - Source Identification
  - Source separation
  - Speaker identification
  - Speaker profiling
  - Speaker verification
  - Speech recognition
  - Speech separation
- Key analysis techniques (selected)

# Speech

## recognition/translation/generation

- Universally use **deep learning architectures** (Take Introduction to Deep Learning 11-785 to understand how these work)

Latest techniques:


- waveform in
  - For ASR: Text out
    - Read <https://arxiv.org/pdf/1609.06773.pdf>
    - Combines both CTC based model and attention based model while training, with improved performance
  - For translation: Translated word sequence out
- Text in
  - For generation (synthesis): waveform out
- Best practice: use available APIs from industry for these tasks, e.g.
  - Google Cloud Speech API
  - IBM Watson Speech to Text
  - IBM Watson Text to Speech
  - Microsoft Azure Bing Speech API
  - Amazon Transcribe
  - Amazon Polly

# Source identification



- Examples of sounds are stored in databases
  - Matching techniques are used to compare signatures
    - Machine learning techniques are used

# Audio event detection



**General Framework for Audio Event Detection**

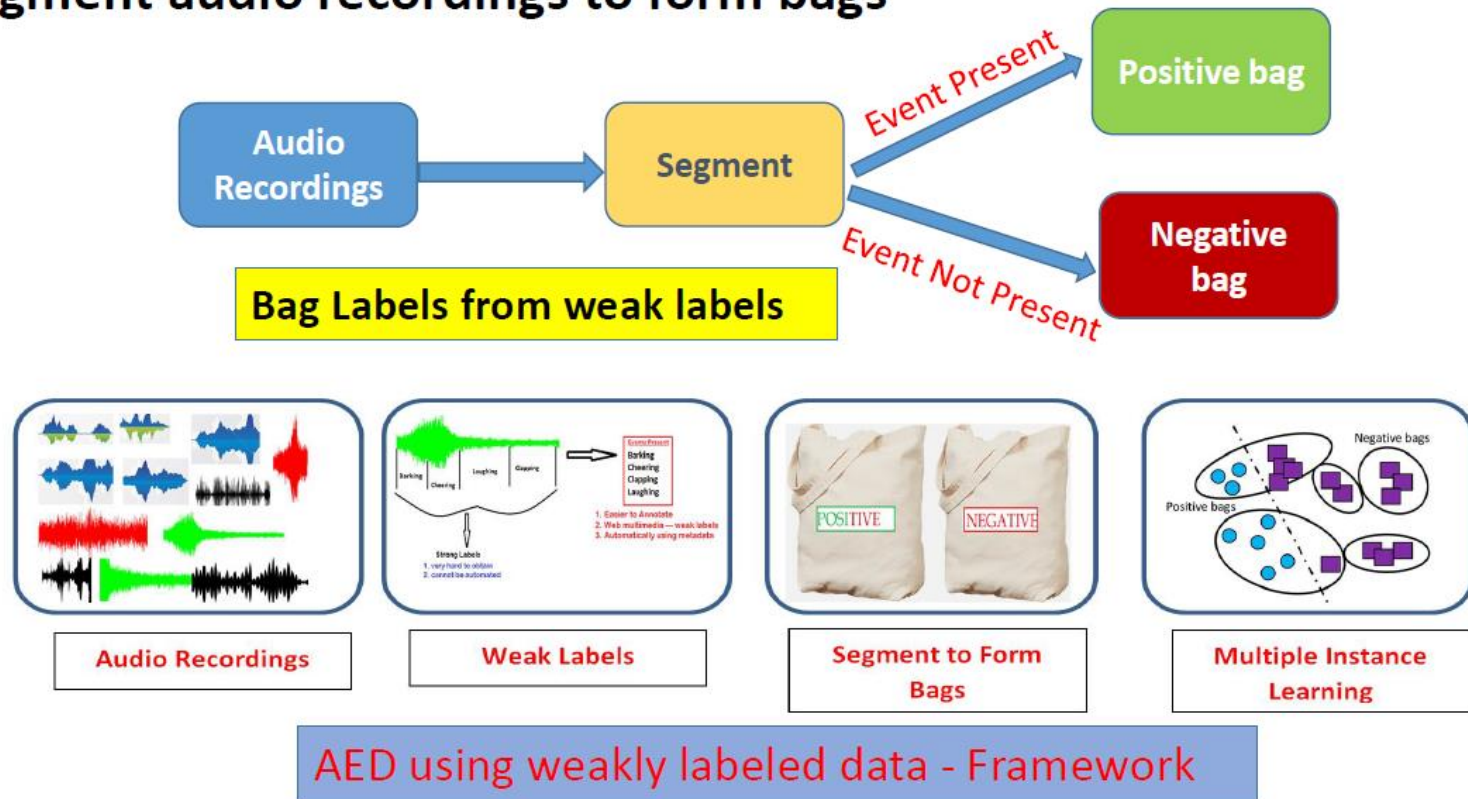
Labeling data with time stamps – Biggest Problem

- AED using strong labels



# Audio event detection

- Segment audio recordings to form bags




- AED using weak labels: Weakly supervised learning
- Machine learning algorithm generally used: **multiple instance learning (MIL)**

# Audio fingerprinting

- “Fingerprints” (patterns) of audio are detected and used to align evidence from multiple audio recordings
  - These may be associated with video recordings


# Audio fingerprinting



- Spectral peaks are aligned

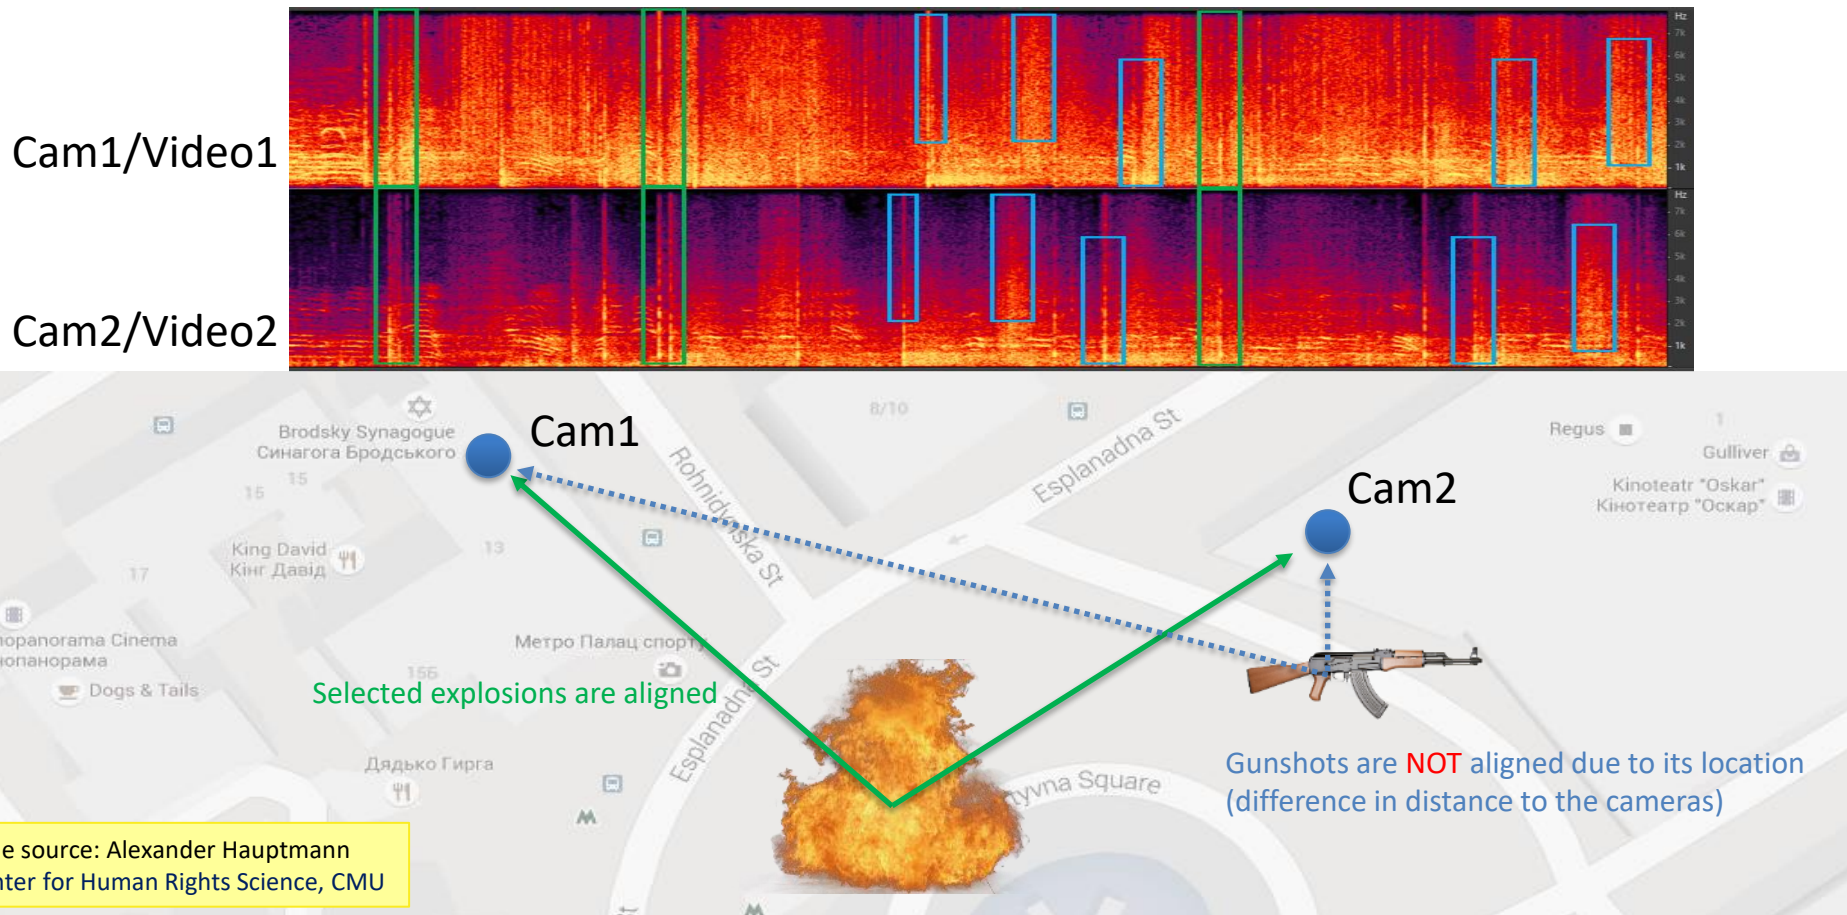


# Audio fingerprinting



# Object positions

- Content-based automatic audio alignment
  - Match video pairs based on broadcast, sirens, explosions, gunshots
- Object positions can be determined by triangulation



# In this lecture

- Digital multimedia: Recording and devices
  - Audio
  - Images
  - Video
  - Text
- Digital multimedia: Processing
  - Audio processing
  - Two generic processing techniques